Group 2 Project Tutorial

William Gomez | Robin Huynh | Clever I Lemus | Luis Rodriguez | Jung Tae Cha

Importing CSV COVID-19 US County JHU & Demographics data into Kibana of Elasticsearch

Step 1 - Download the data from Kaggle

1. You need to launch and open Kibana
2. We are going to download the data files from the Kaggle site (https://www.kaggle.com/datasets/headsortails/covid19-us-county-jhu-data-demographics) and then import it into Elasticsearch.

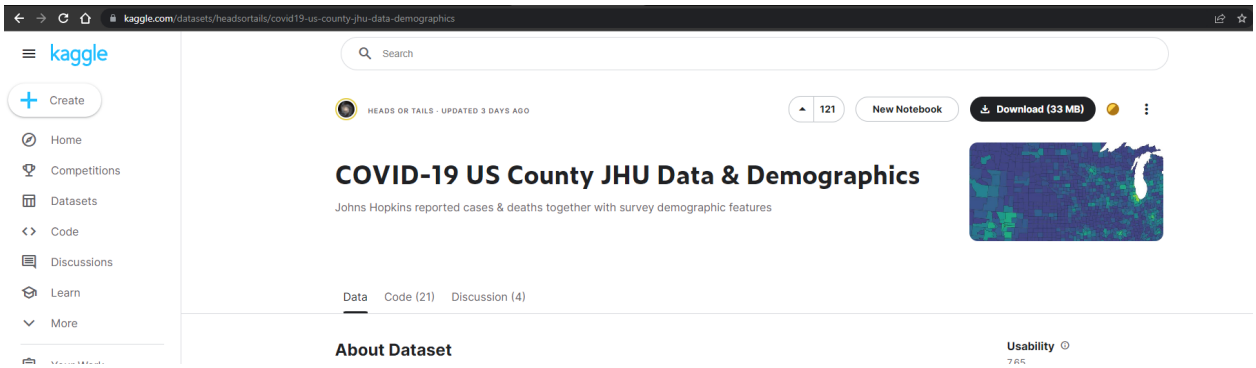

3. Once you get into the website, scroll down and click on the download button:

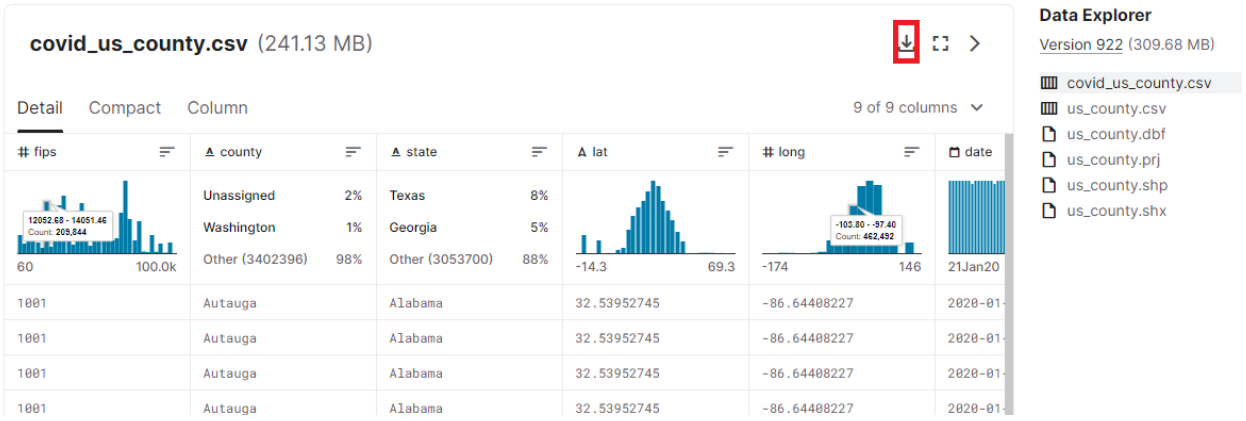

4. The file will be downloaded in a .ZIP format, let's extract it to our desired location.

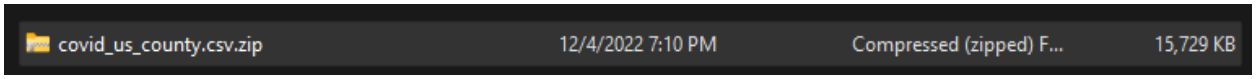

Step 2 - Import the data from CSV and Creating Data View

Now that we have the data to analyze, we need to import it and set up an index.

1. Within Kibana, click on **Upload a file**.

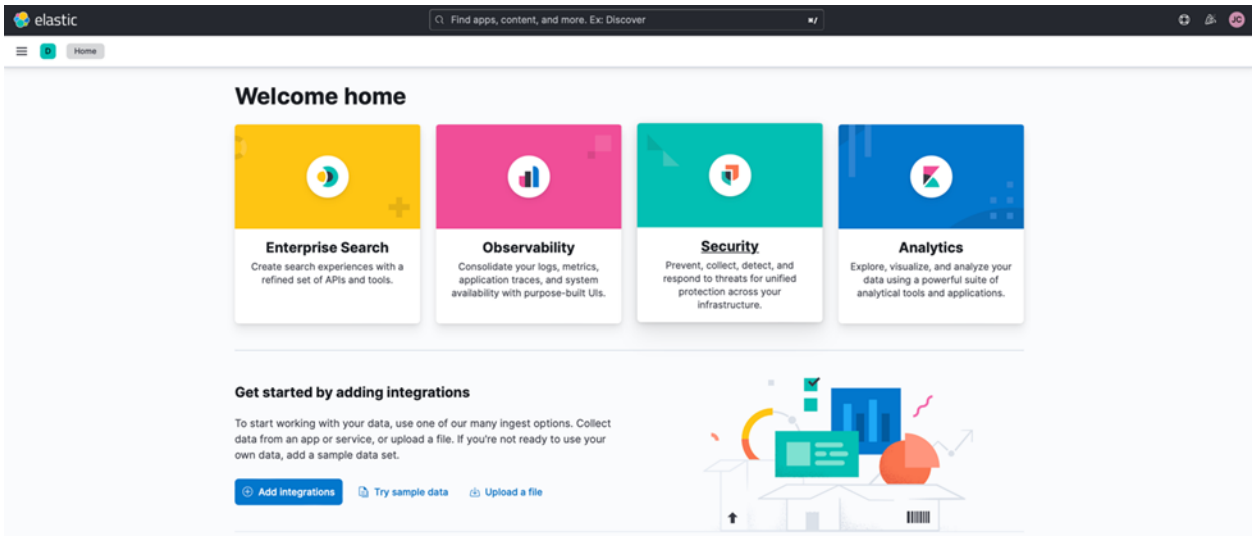

2. Let's select our CSV data set by clicking on "**Select File**".

*** If you get an error telling you that there is a limit to the upload size, we can fix that.

a. Let's click on the menu button.

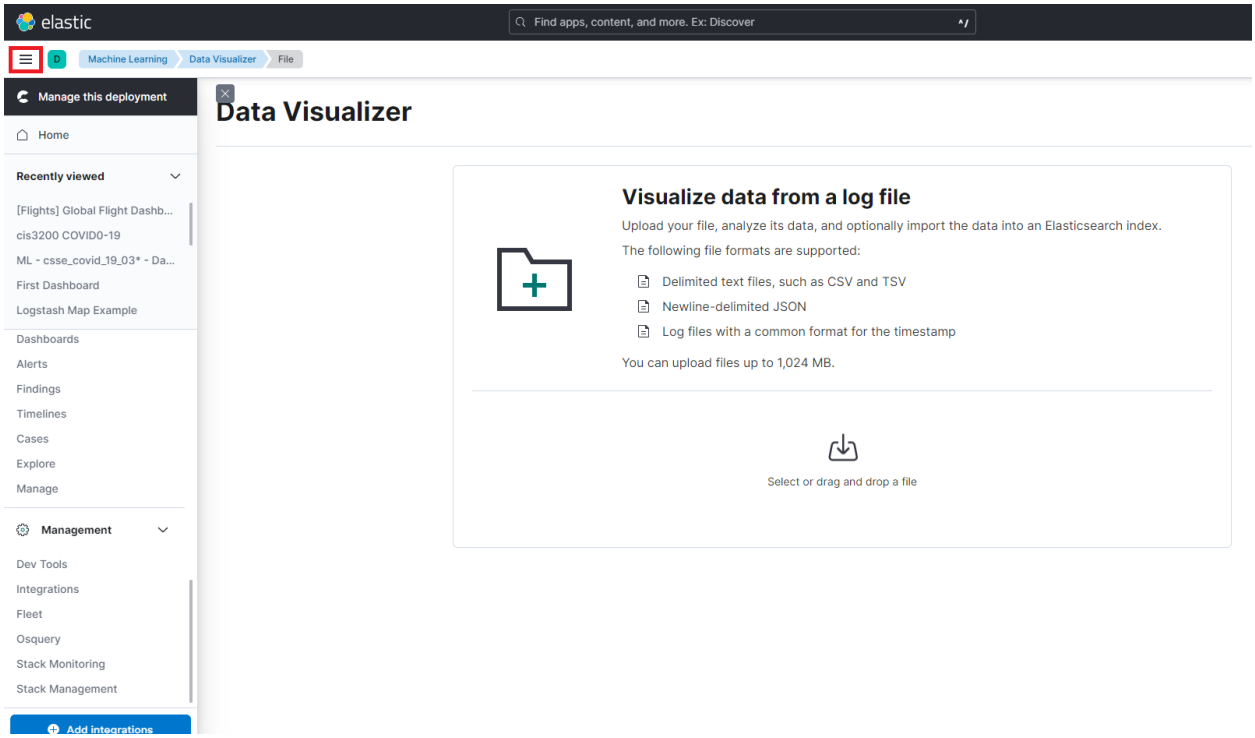

b. Let's select "**Stack Management**" under the **"Management**" subnav.

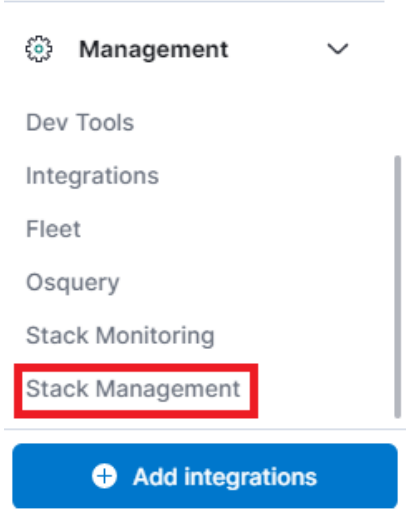

c. Let's click on "Advanced Settings"

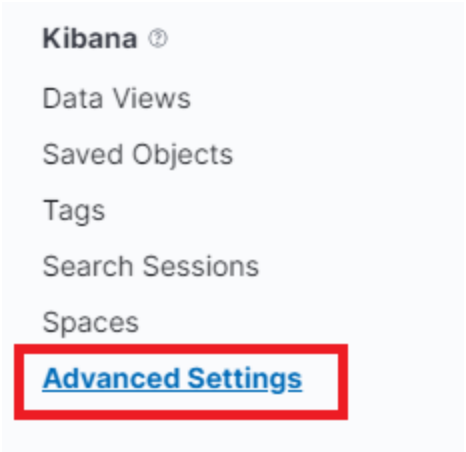

d. Under "**General**" settings, we scroll down to **"Maximum file upload size**". It seems like the default is 100MB, we can change it whatever size we desire, with the maximum upload size being **1GB**, so let's change it to that.

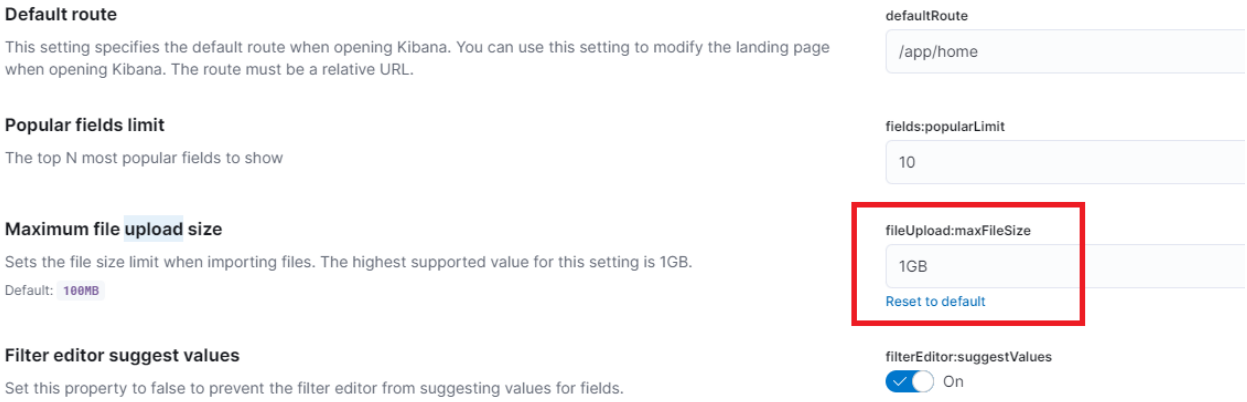

With this, we fixed our upload issue and can continue uploading our data to Kibana.***

4. We select the covid_us_county.csv file, shortly it will show a preview of the data.

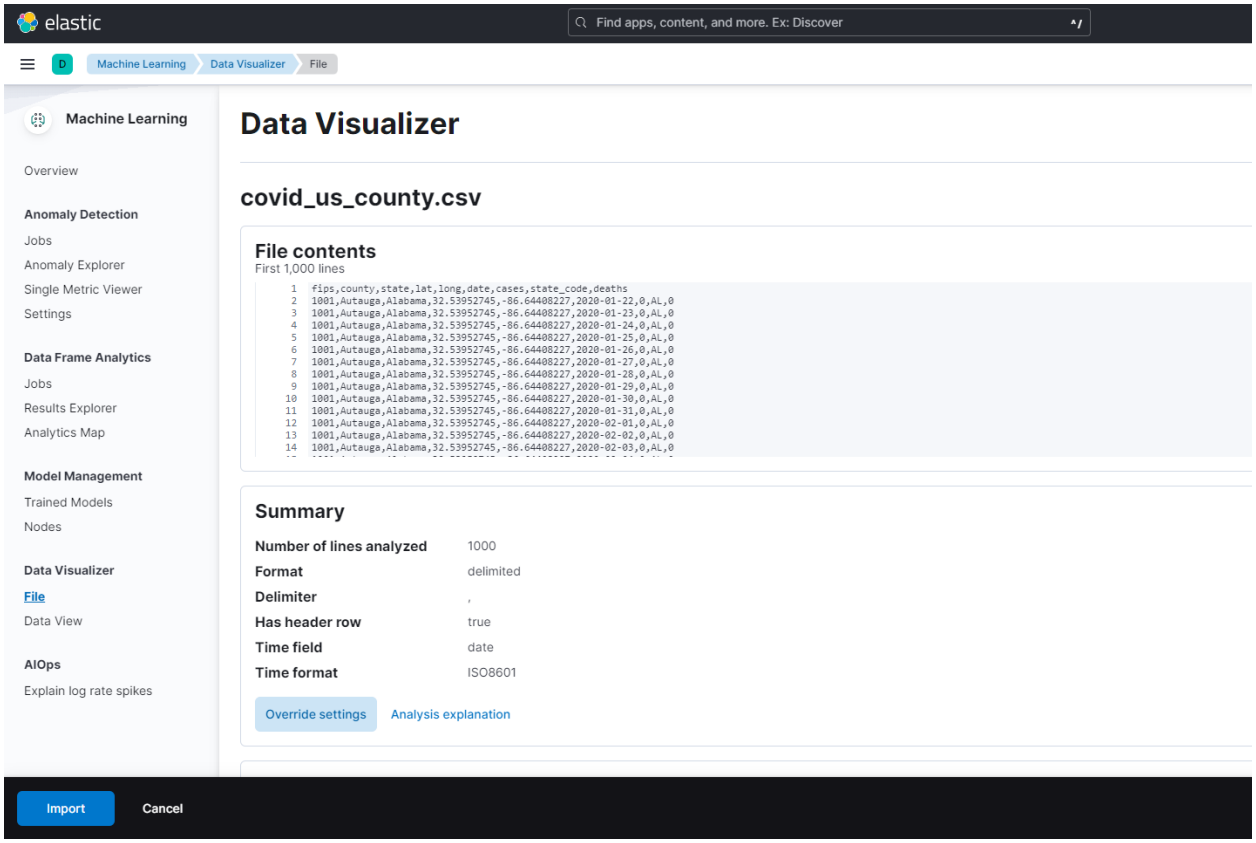

5. Now let's select **Override settings** and change the column names as: "Lat" to "**Latitude**" and "Long_" to "**Longitude**". Then, we click on **Apply**.

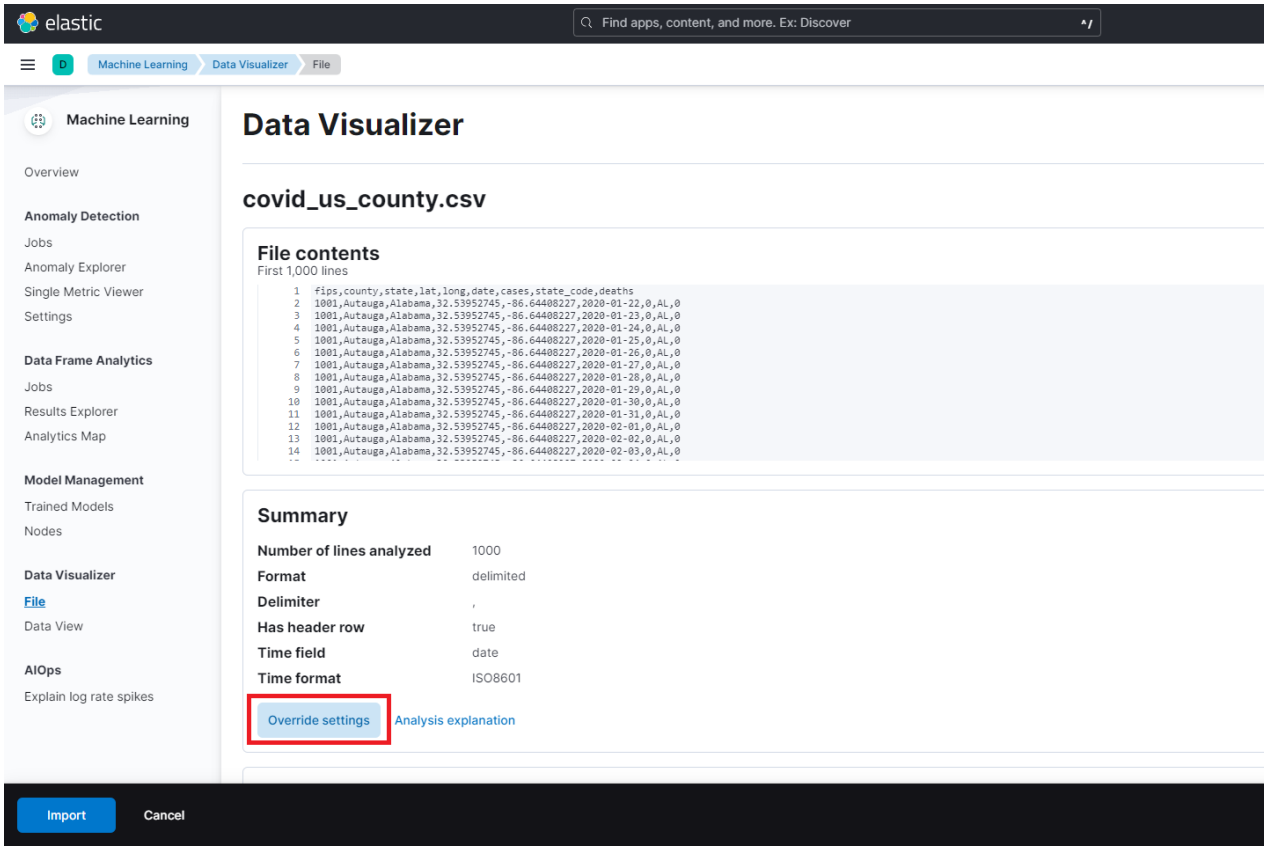
elastic
Find apps, content, and more. Ex: Discover
D
Machine Learning
Data Visualizer
File
Machine Learning
Overview
Anomaly Detection
Jobs
Anomaly Explorer
Single Metric Viewer
Settings
Data Frame Analytics
Jobs
Results Explorer
Analytics Map
Model Management
Trained Models
Nodes
Data Visualizer
File
Data View
AIOps
Explain log rate spikes
Data Visualizer
covid_us_county.csv
File contents
First 1,000 lines
fips,county,state,lat,long,date,cases,state_code,deaths
1001,Autauga,Alabama,32.53952745,-86.64408227,2020-01-22,0,AL,0
1001,Autauga,Alabama,32.53952745,-86.64408227,2020-01-23,0,AL,0
1001,Autauga,Alabama,32.53952745,-86.64408227,2020-01-24,0,AL,0
1001,Autauga,Alabama,32.53952745,-86.64408227,2020-01-25,0,AL,0
1001,Autauga,Alabama,32.53952745,-86.64408227,2020-01-26,0,AL,0
1001,Autauga,Alabama,32.53952745,-86.64408227,2020-01-27,0,AL,0
1001,Autauga,Alabama,32.53952745,-86.64408227,2020-01-28,0,AL,0
1001,Autauga,Alabama,32.53952745,-86.64408227,2020-01-29,0,AL,0
1001,Autauga,Alabama,32.53952745,-86.64408227,2020-01-30,0,AL,0
1001,Autauga,Alabama,32.53952745,-86.64408227,2020-01-31,0,AL,0
1001,Autauga,Alabama,32.53952745,-86.64408227,2020-02-01,0,AL,0
1001,Autauga,Alabama,32.53952745,-86.64408227,2020-02-02,0,AL,0
1001,Autauga,Alabama,32.53952745,-86.64408227,2020-02-03,0,AL,0
Summary
Number of lines analyzed 1000
Format delimited
Delimiter ,
Has header row true
Time field date
Time format ISO8601
Override settings
Analysis explanation
Import
Cancel

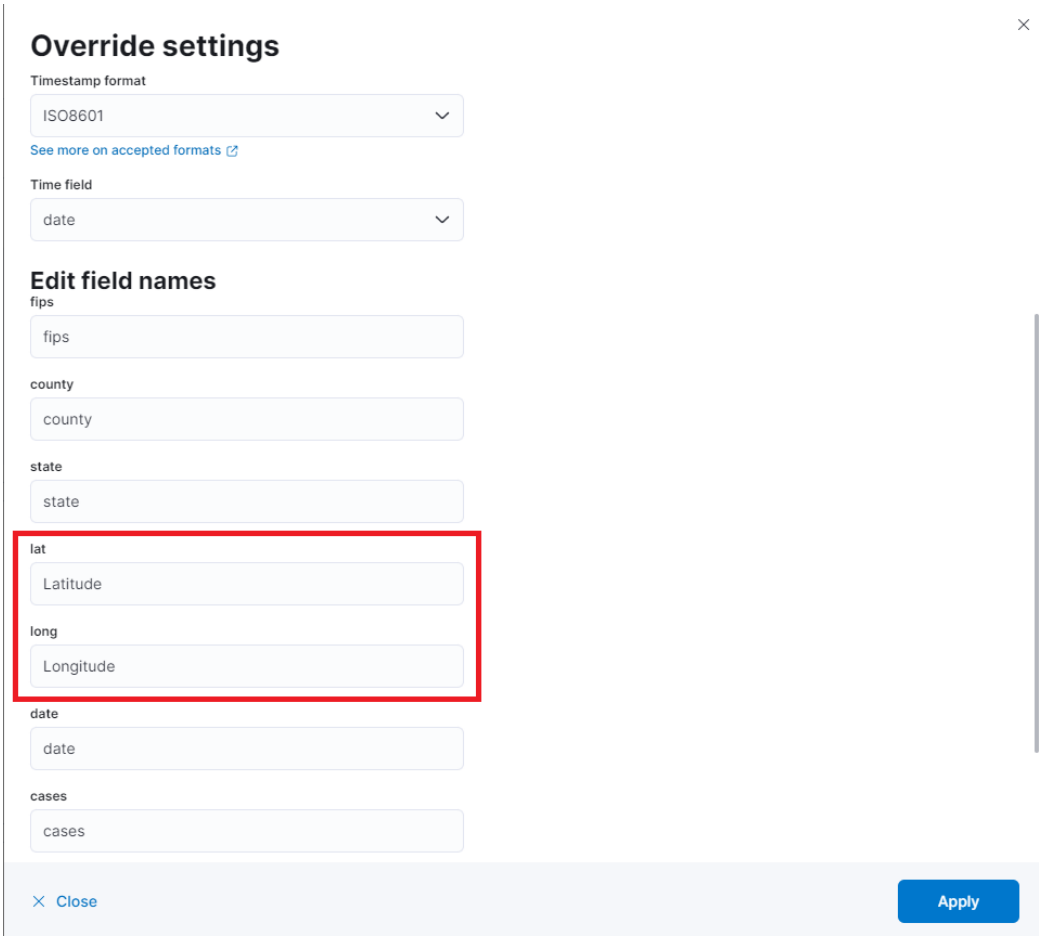

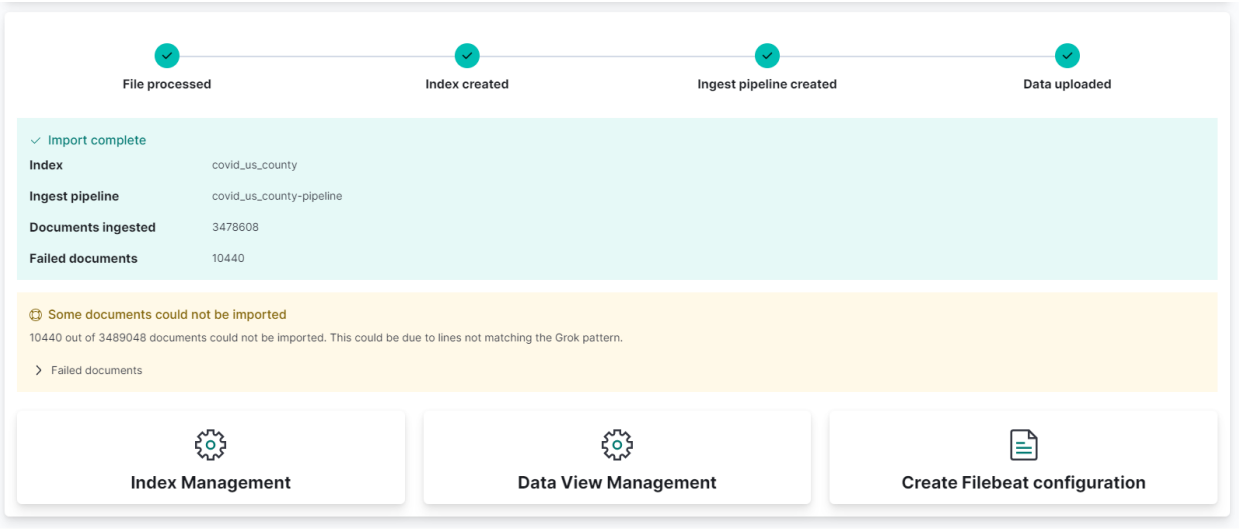

6. Now that we have our data uploaded, we can create our data view which will allow us to use that for many other things, such as visualizations and machine learning. We're going to click on Data Views under Kibana.

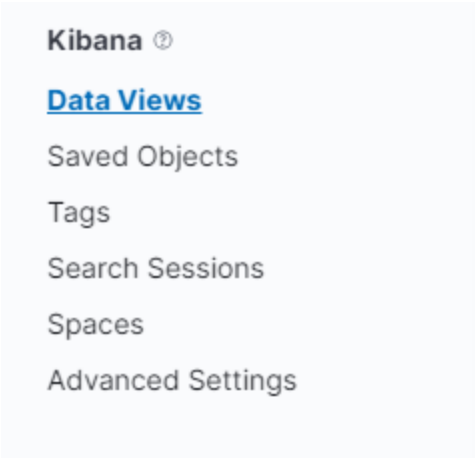

Select “Create Data View”.

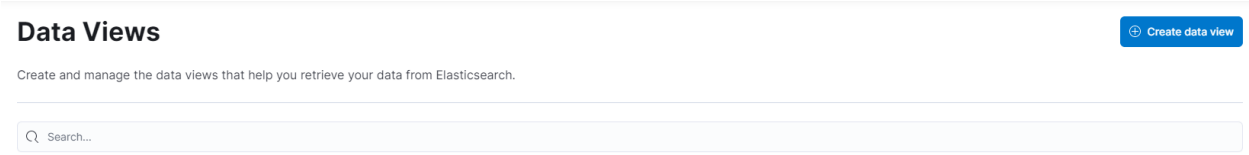

Let’s choose an index pattern name and let’s make sure that it matches our uploaded data. Also, let’s make sure that timestamp is the selected timestamp pattern.

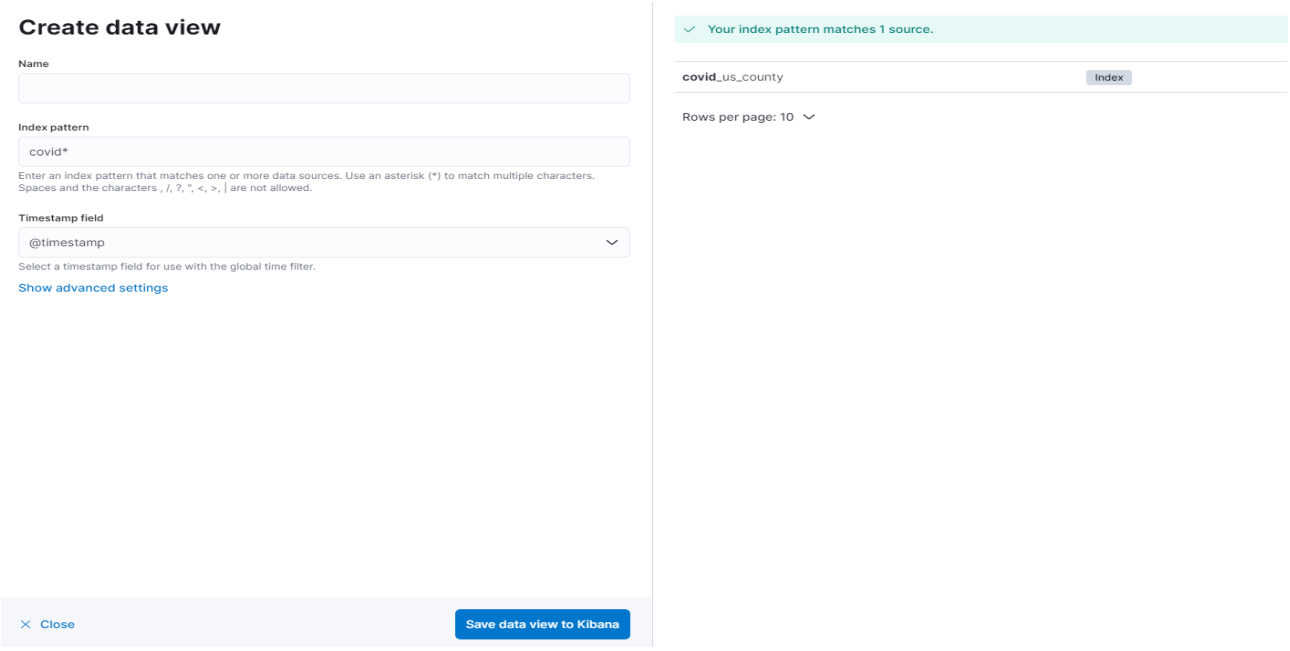

Creating a Map on Elasticsearch

1. Now that we have our data view created. We can do many things with it on ElasticSearch. Let's start by creating a map using our geolocation that we created in the earlier steps. To do so we will head over to "**Maps**" under **analytics**.

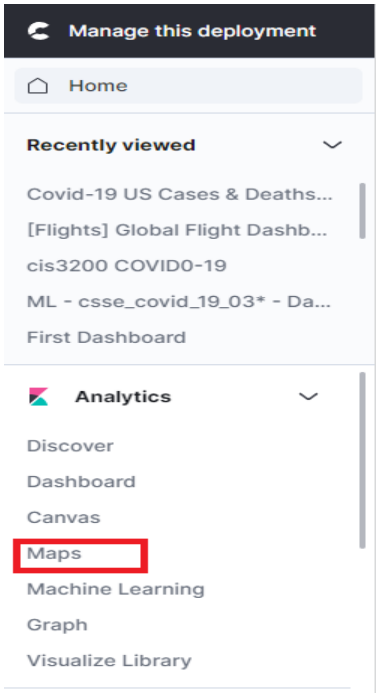

2. We are going to click on "**Create Map**".

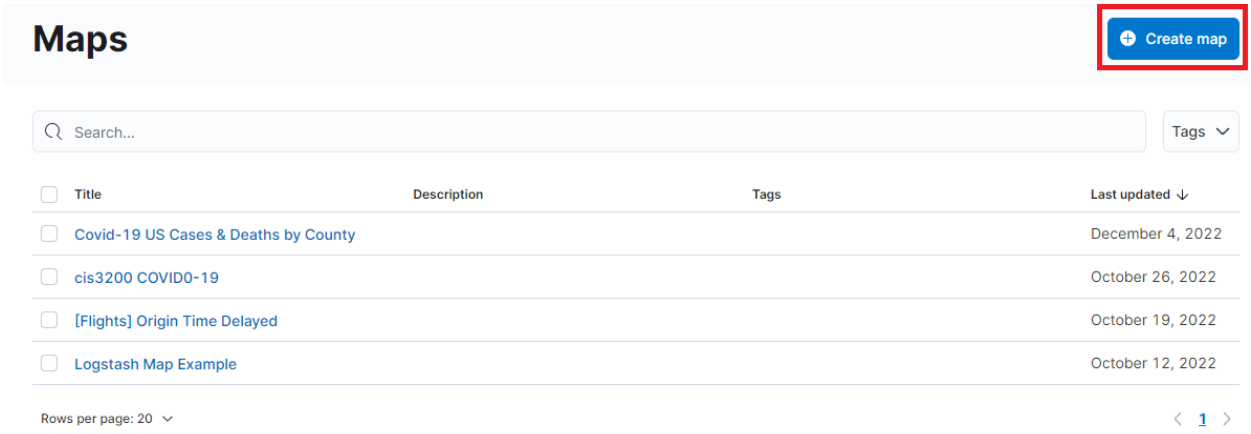

3. Let’s click on “**Add Layer**” and let’s select “**Clusters**”

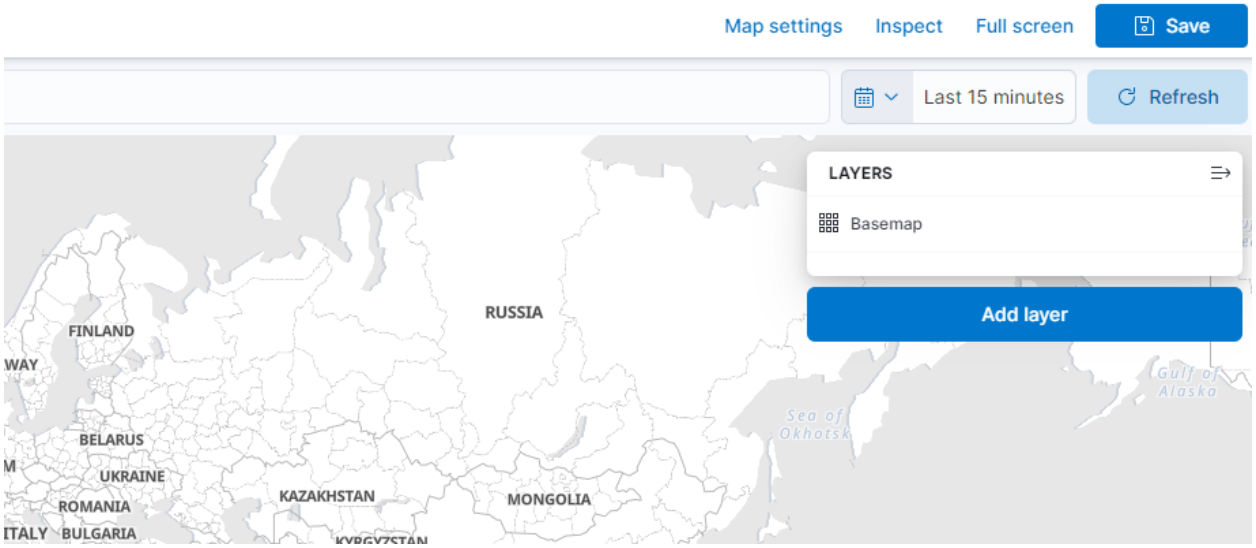

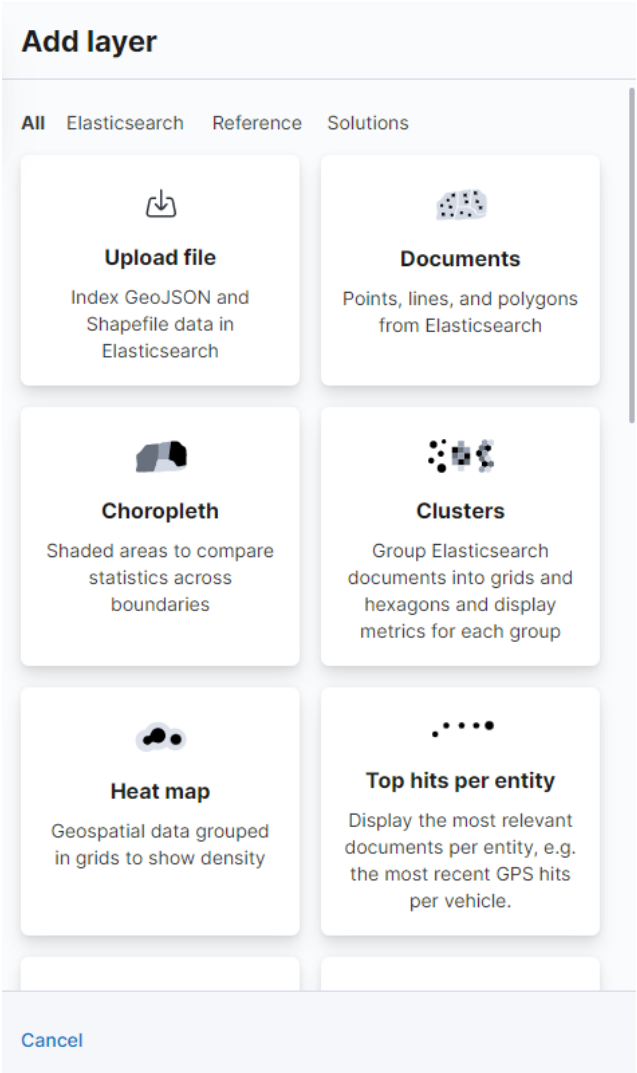

4. Let's select the data set we uploaded. Let's also make sure we select the desired time range. For our dataset it will be **Jan 22,2021** to **Dec 3, 2022**. Time shouldn't matter.

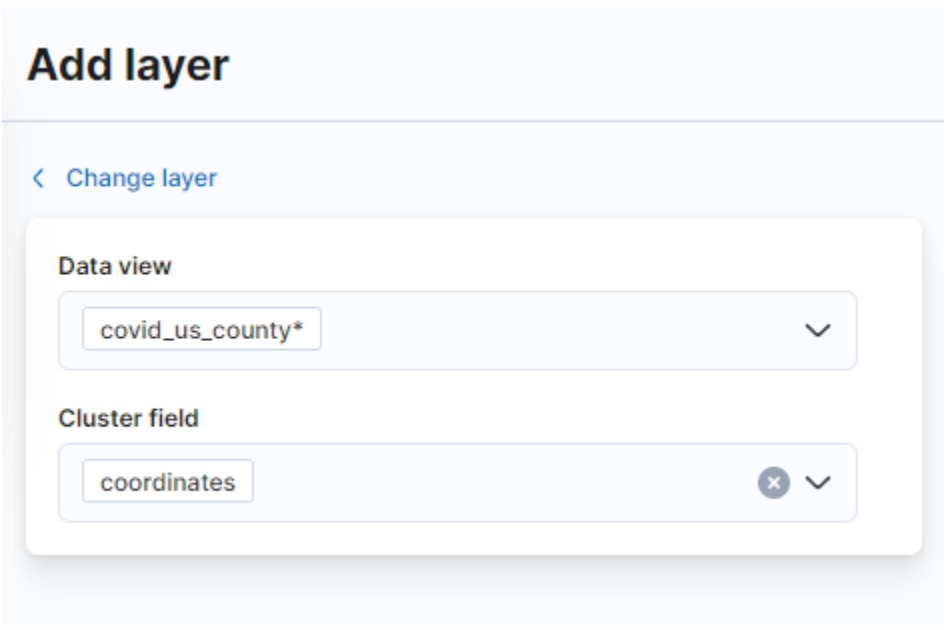

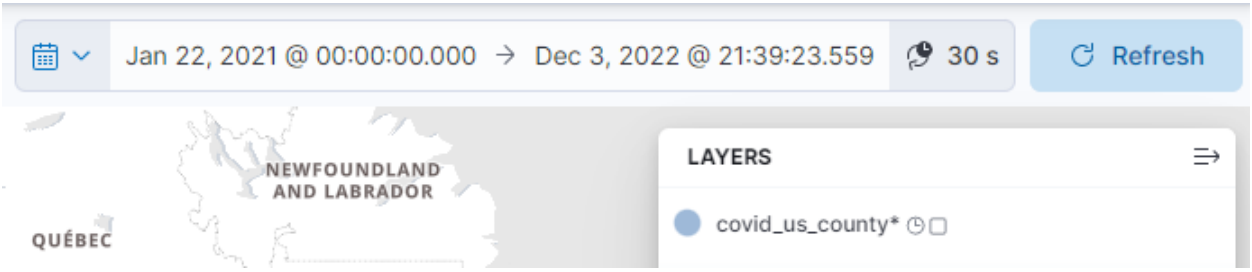

5. Once we have selected the time range and selected the coordinates field with a custom label name, we are greeted with our map. We can click save so that we can use it in the future for any possible need.

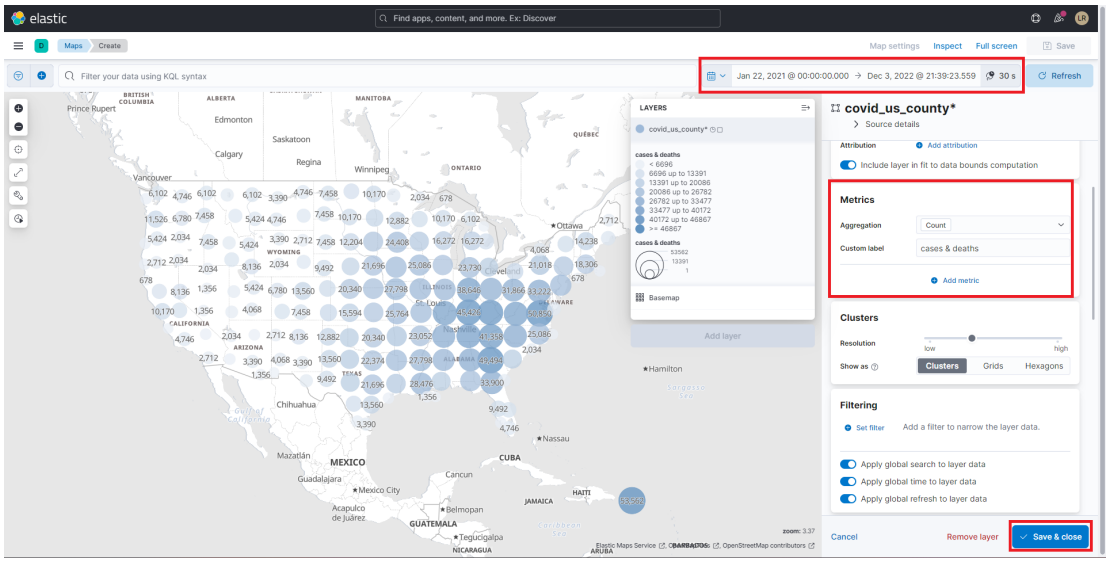

Creating a Pie Chart

1. At this step, we already have our data set uploaded and we already created our data set. If not, please refer to the beginning of this document.
2. On the left hand side, let's open up the sliding elastic menu and let's click on "Visualize Library" under "Analytics".

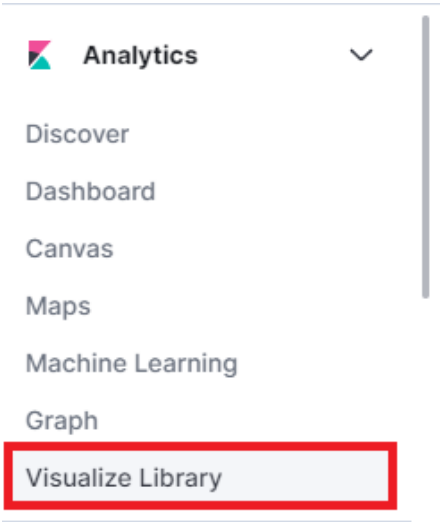

3. You will see the following screen, it should be blank if this is your first visualization. We'll simply click on the blue button that says "Create Visualization".

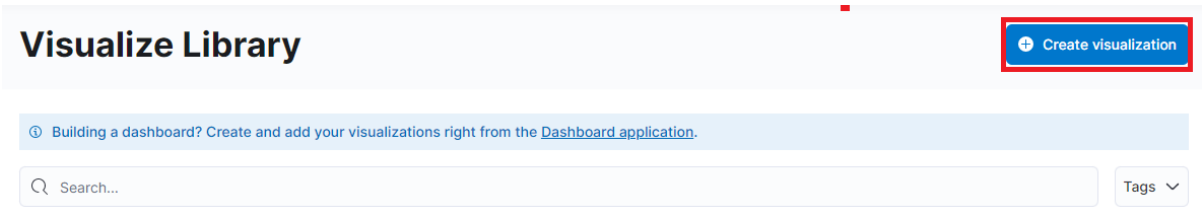

4. The page will prompt us with several options that elastic provides for our data, let's click on "Explore Options" under the "Aggregation based" section.

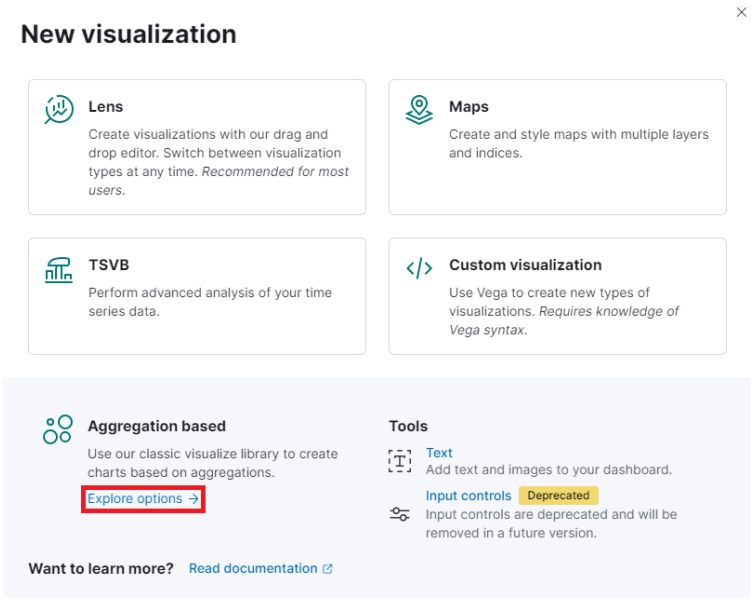

5. Let's select our desired data set.

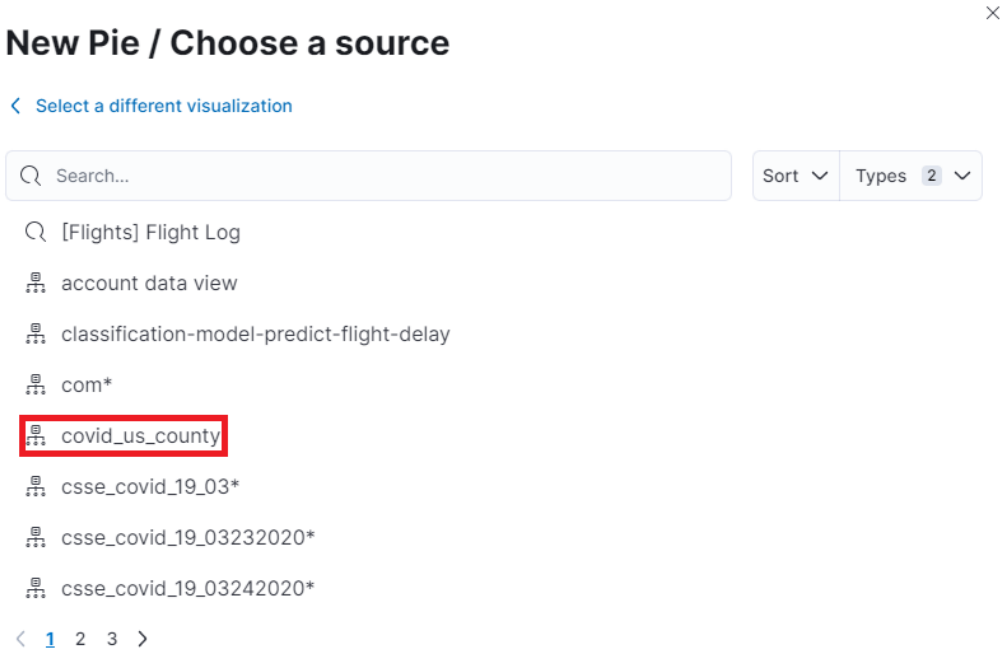

6. The following step requires us to select the date range from our data set. We must also use the following options. For aggregation we want "**Top Hit**". For the field we will be using "**cases**" since that's what we're looking for in this situation. Under "**Aggregate with**" we went with **Sum**, but Elastic provides other options such as "Average", "Max", or "Min", it depends on what it is you are trying to present with your visualization.

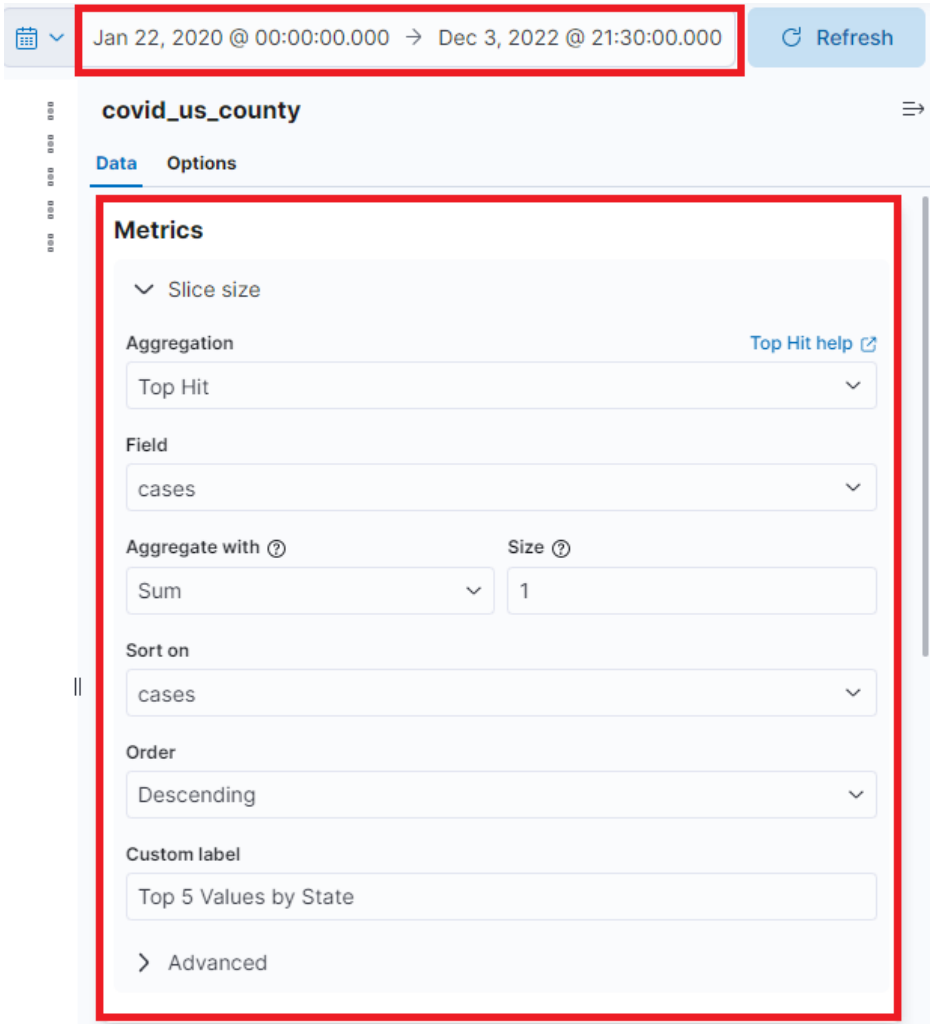

7. Now that we have customized all of our desired fields, let's click on update so that Elastic can apply our changes.

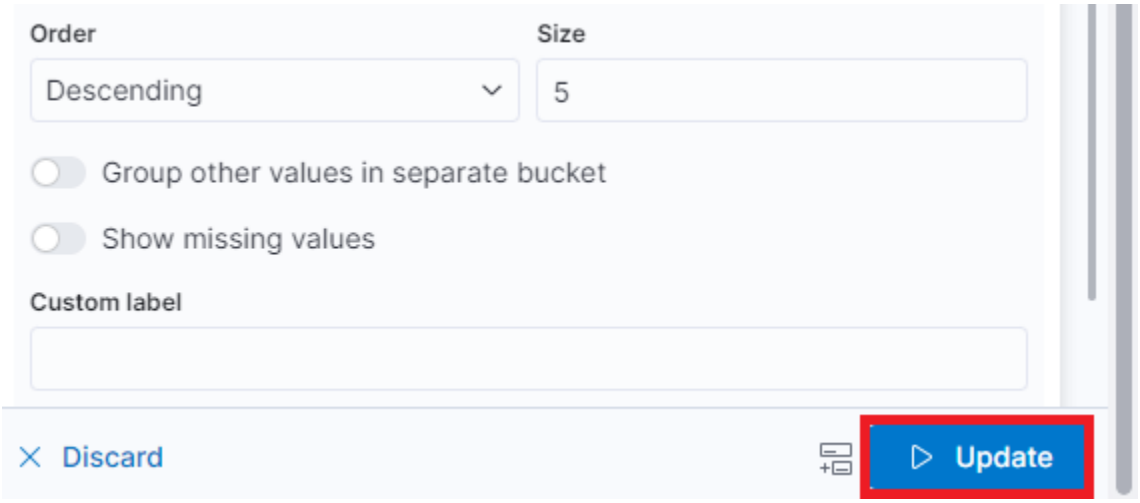

8. As you can see, we should now be able to see the top 5 cases sorted by state.

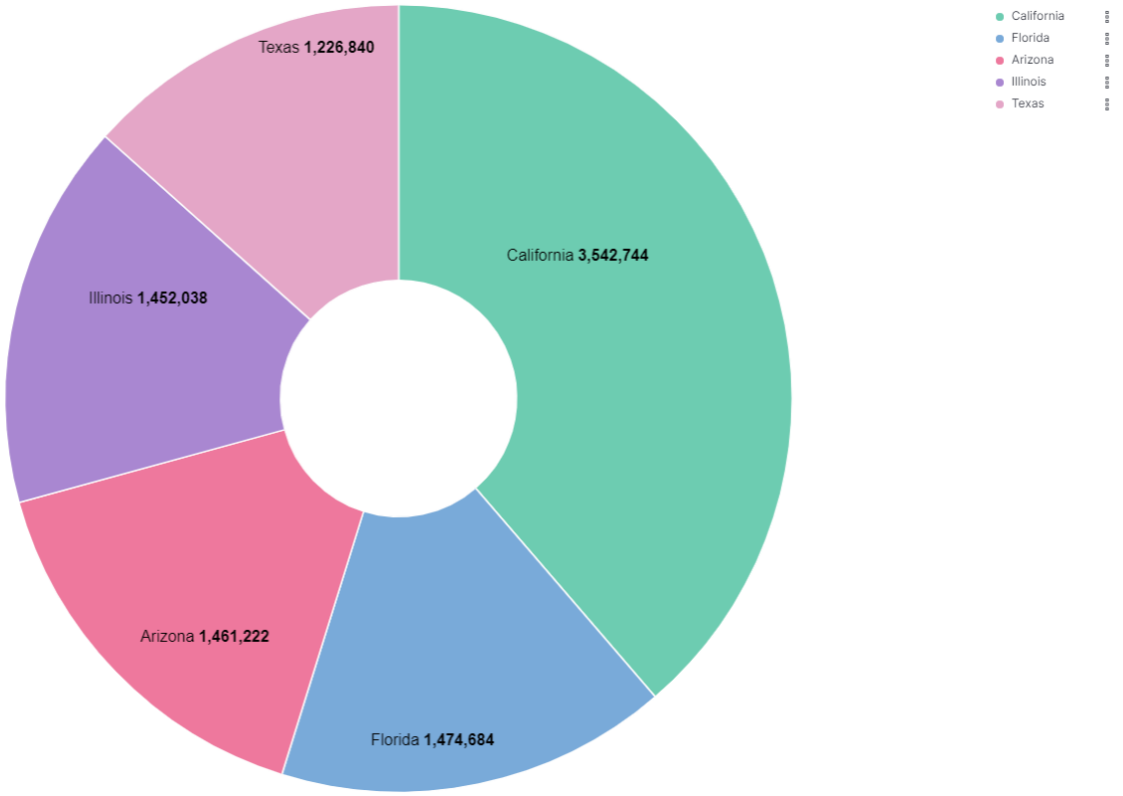

9. In instances, you may be asked to include the timestamp that was created back when we were creating our data view. We can include this in our pie chart, with a couple of extra steps. Let's start by adding an additional bucket.

10. Select “**Split slices**”

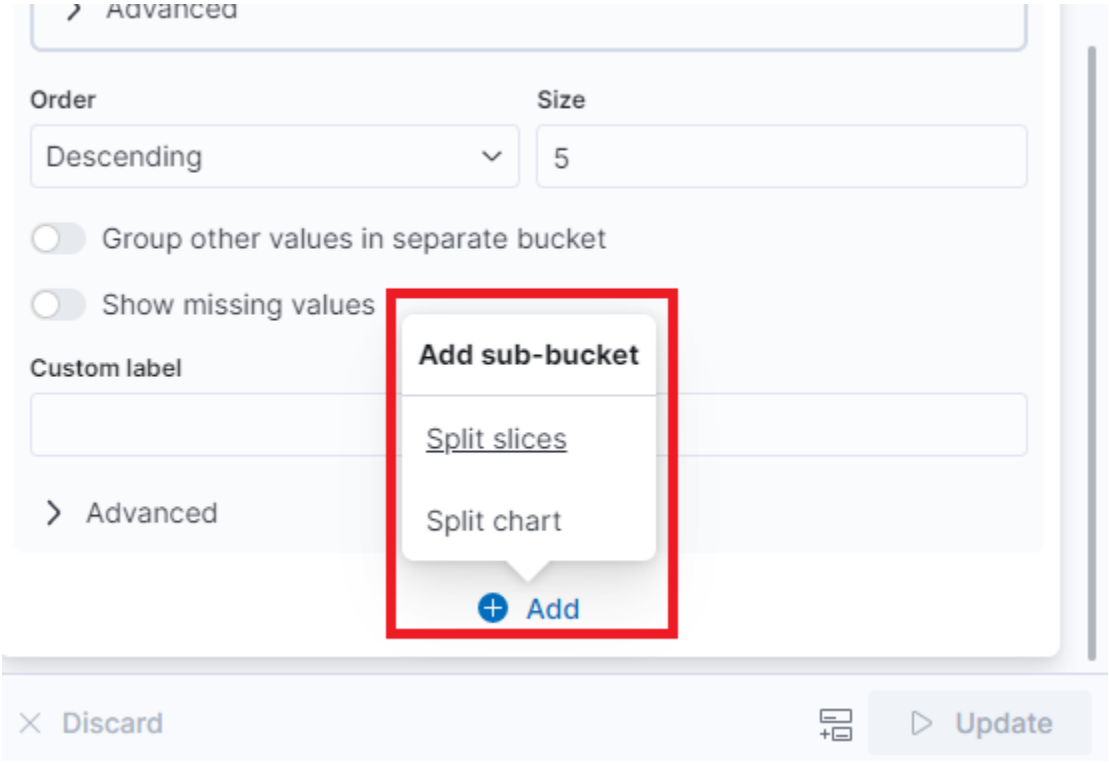

11. Let’s select “**Data Range**” for our Sub aggregation and choose “**timestamp**” for the field. Of course, we must also select the actual dates from our data set.

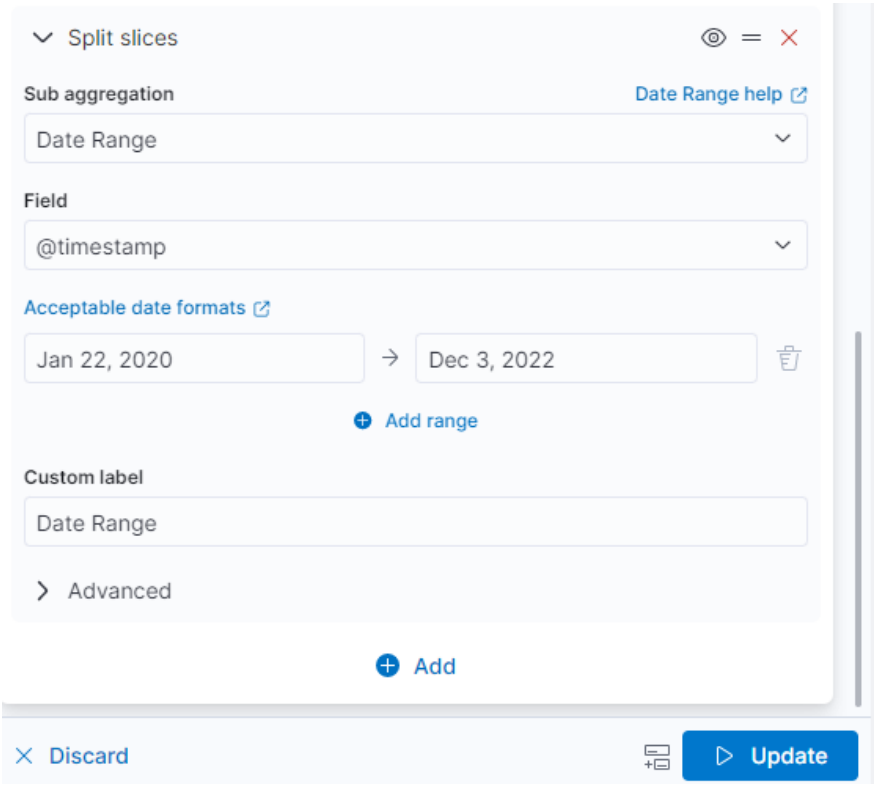

12. We will click through advance and let it know we want 5 cases and for it to be in Descending order.

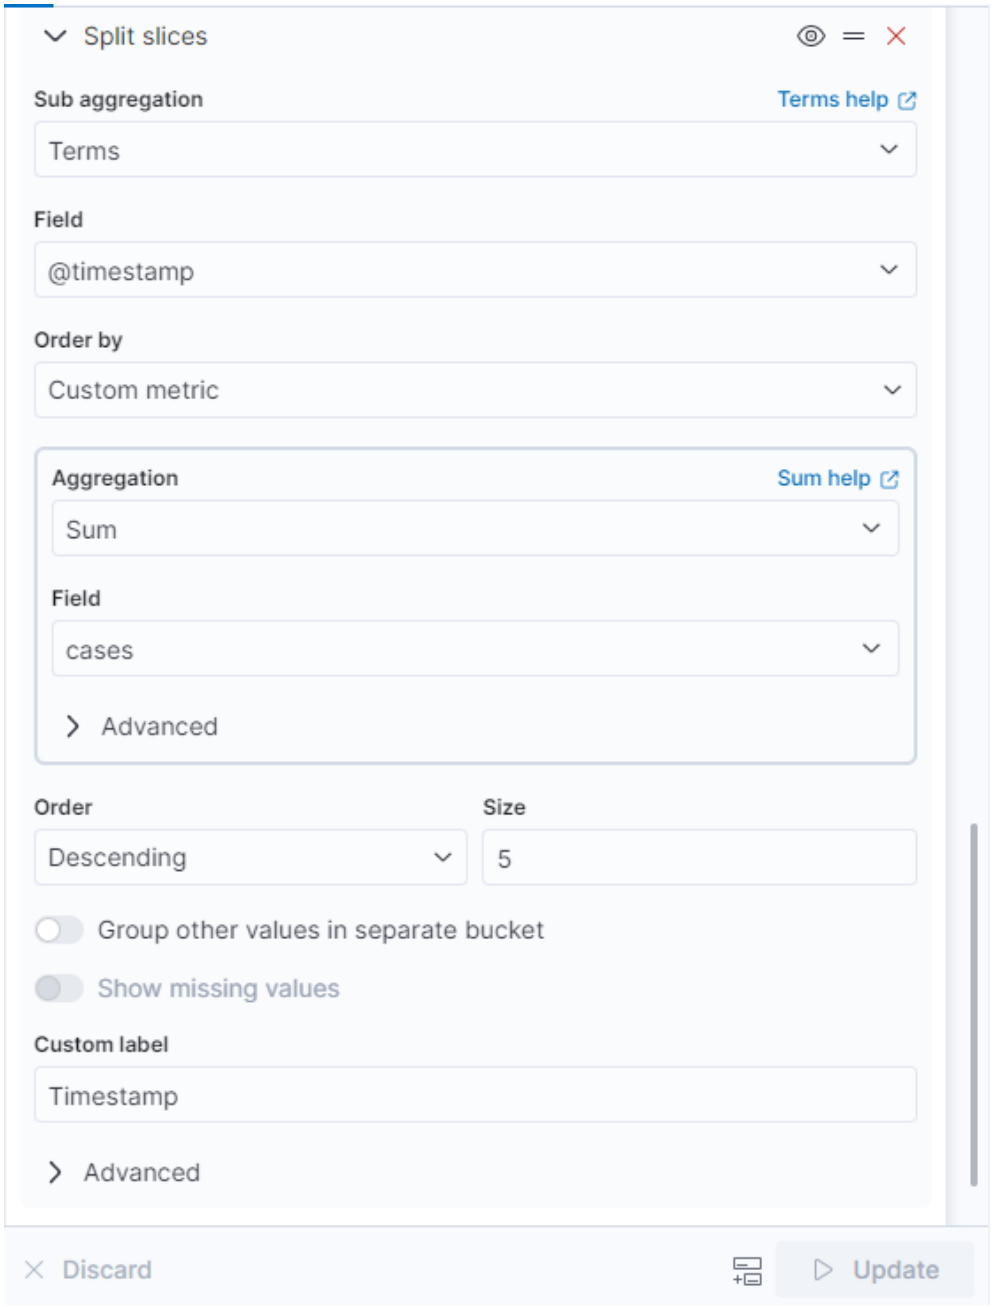

13. Here are the results, a new pie surrounds our previous pie, this time, it provides the date and shows the number of cases updated on that date.

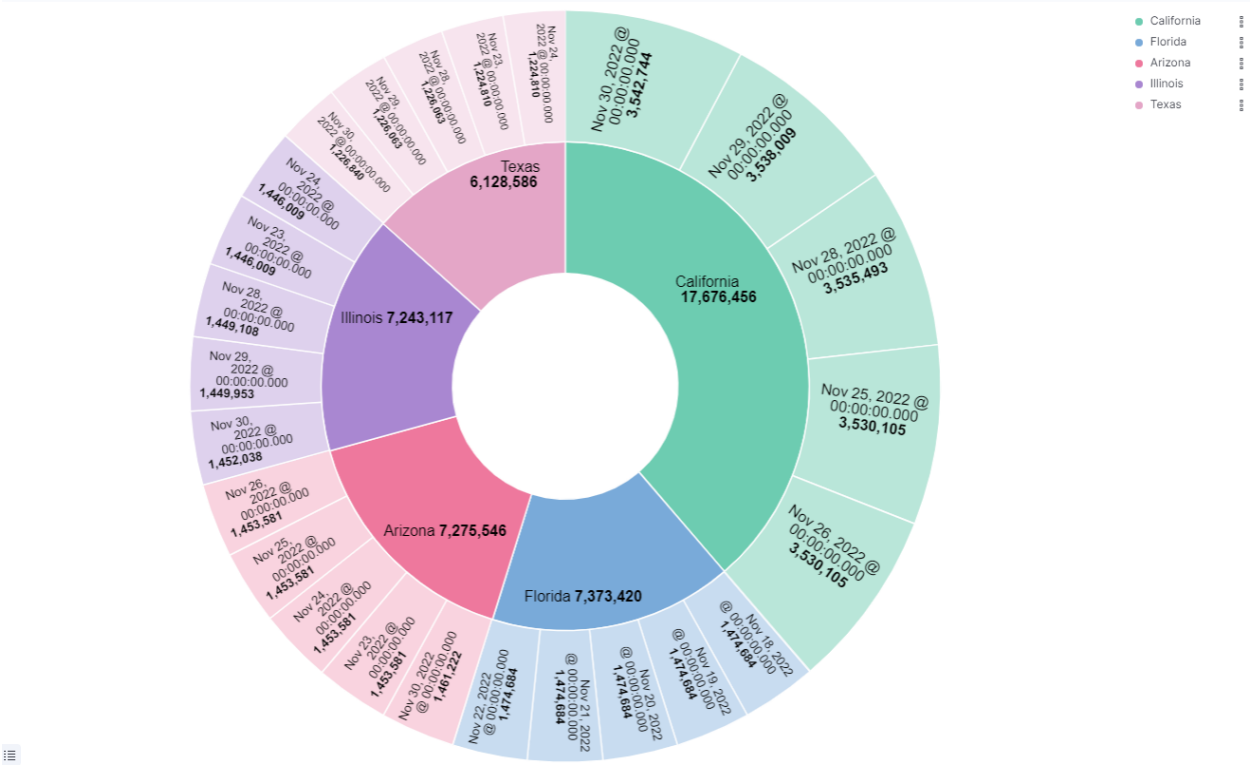

Creating a Regression Job in Elastic's Machine Learning

Another feature that Elastic provides is several analytics that use machine learning. We can identify anomalies on data. We can also make predictions based on the data we feed it. That's what we will be doing in this situation. Our data is based on COVID-19 cases based on county, so let's use Elastic's Machine Learning to predict future cases, per county.

1. First, let's click open the sliding menu on Elastic, scroll down to "**Analytics**" and select "**Machine Learning**".

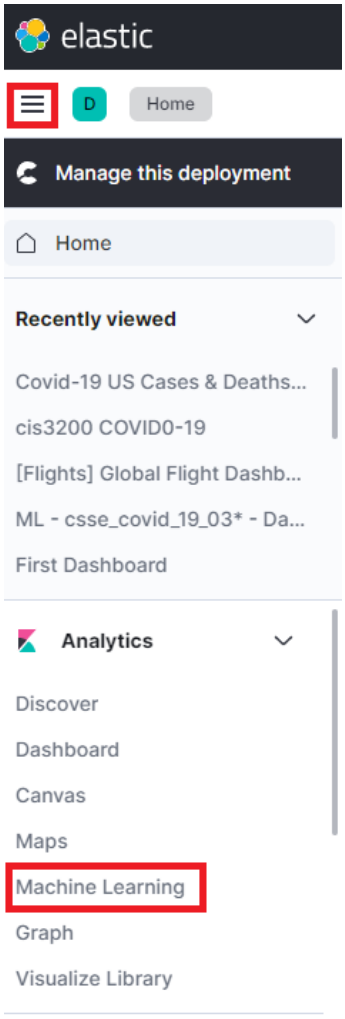

2. We are first greeted with the Overview page. This is where we will be able to see all the jobs we create in the future. For now, on the left hand side, under **"Data Frame Analytics**", lets click on "**Jobs**".

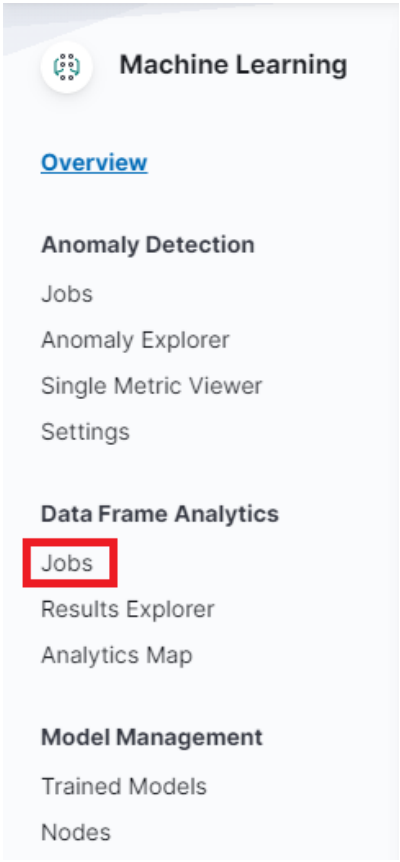

3. The page shows us all of our Data Frame Analytics jobs, we're going to click on the blue button that says "**Create Job**".
4. On this page, we can see all of our data views. Let's choose the data view that we created for this project. We titled it "**covid_us_county**".

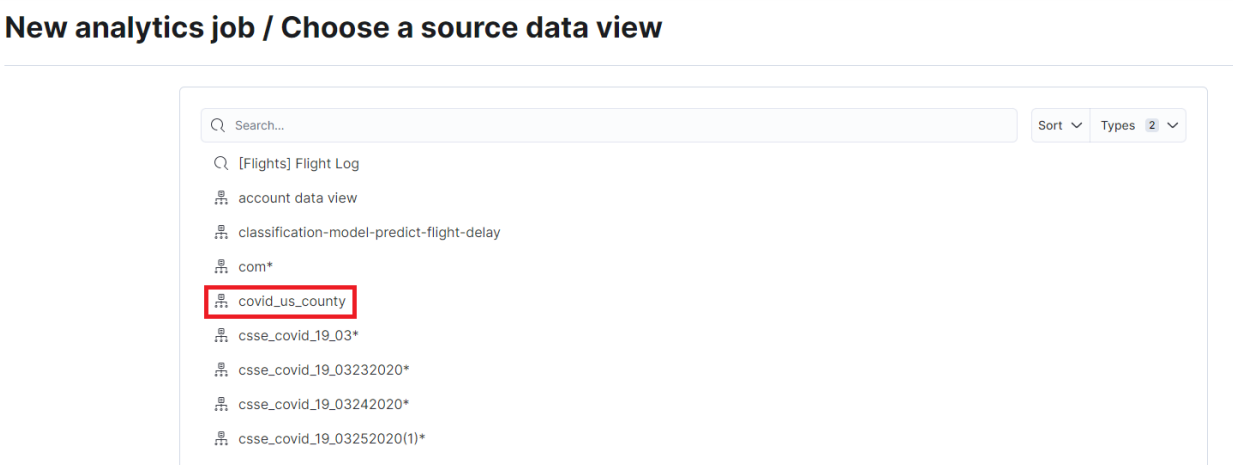

5. After selecting the data view, Elastic wants to know what we would like to do with the data. We're gonna be predicting numerical values in the data set, which is exactly what “Regression” does, so let’s choose that option. Since we’re predicting cases, let’s make sure that for the query section we choose **“cases > 1**”. We also want to set our “**Dependent Variable**” to be “**cases**”.

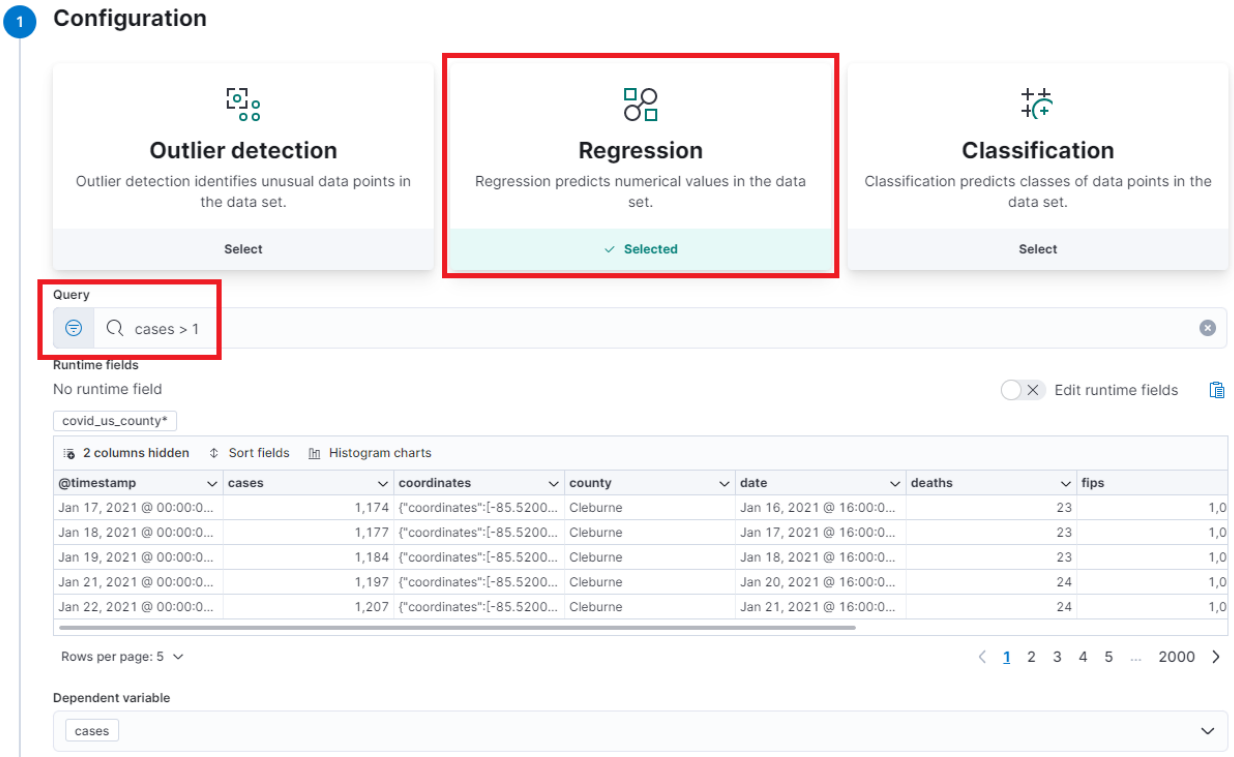

6. To start the process of machine learning, we’re gonna go through 4 more steps to set up the job. Please refer to the following images so that we can set up our job correctly.
    a. For our fields, since we are trying to predict per county, let’s make sure that the field is selected.

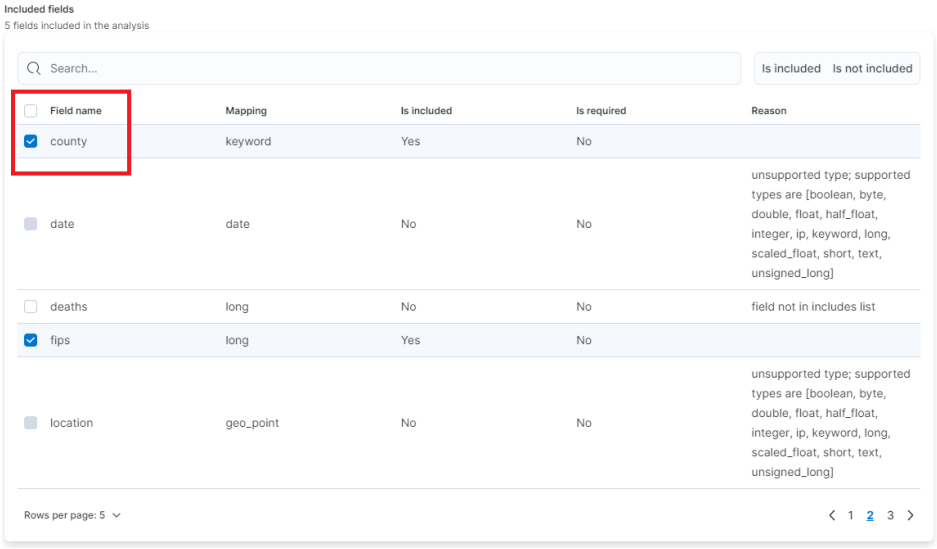

b. For our training percent, we want our prediction to be as accurate as possible, so let's choose 90 for the Training percent. Keep in mind that the higher this percent is, the LONGER it will take to show the results.

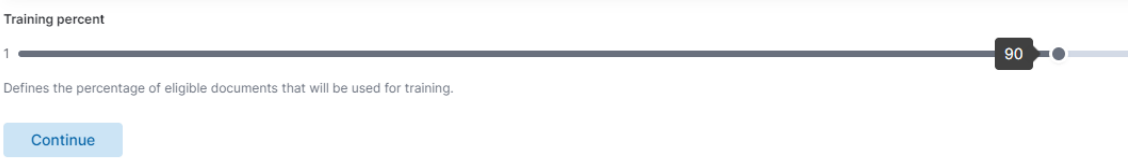

c. Let's set 3 for our feature importance value and let's keep the recommended memory limit.

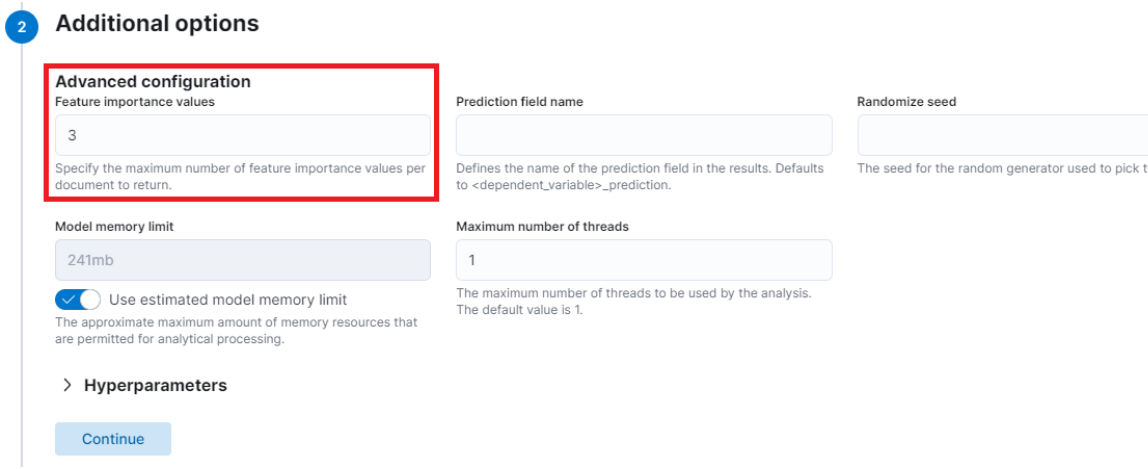

d. Let's name our job.

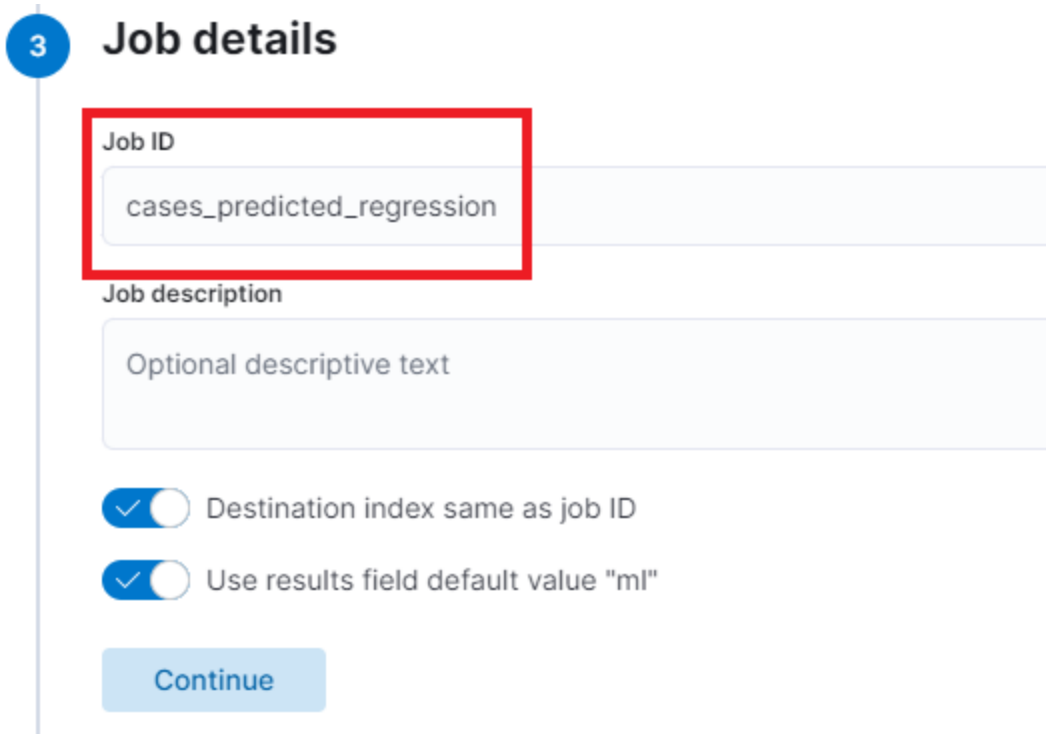

e. The validation step is just telling us that we have our dependent variable and it's letting us know that by choosing to have a high training percent and a feature importance, the process may take longer. It's informational, we can click on continue.

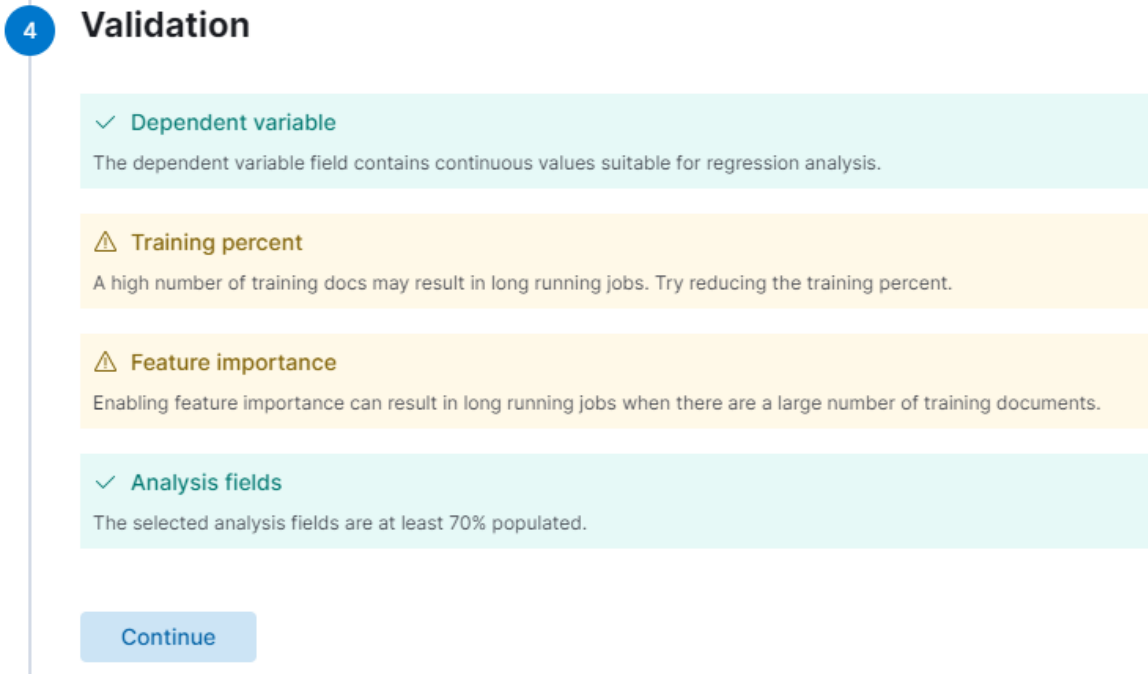

f. We're done setting our specifications. Let's click on the "**Create**" button and now the fun part begins: Waiting!

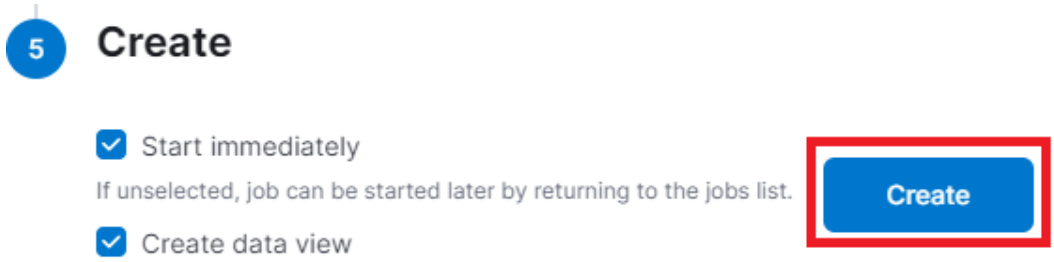

7. Now that progress is complete, let's take a look at our results. Your results should be different than ours when it comes to your Mean Square error (MSE). When you get a 1, it means that the prediction has perfect accuracy, which is highly unlikely. Due to time constraints, we were unable to rerun the machine learning process to get more accurate results.

Results

Total docs
>10000

Showing documents for which predictions exist

12 columns hidden | Sort fields | Histogram charts

ml.is_training	ml.probable cases_pr...	probable cases	ml.feature_importance
false	127,293.573	129,474	[{"feature_name":["cases...
false	131,317.957	132,834	[{"feature_name":["cases...
false	132,559.508	136,771	[{"feature_name":["cases...
false	135,970.66	138,625	[{"feature_name":["cases...
false	139,857.324	141,290	[{"feature_name":["cases...
false	142,350.356	141,877	[{"feature_name":["cases...
false	142,350.356	143,805	[{"feature_name":["cases...
false	147,781.788	146,990	[{"feature_name":["cases...
false	155,976.072	158,967	[{"feature_name":["cases...
false	163,483.668	161,284	[{"feature_name":["cases...
false	172,048.912	169,308	[{"feature_name":["cases...
false	172,048.912	171,050	[{"feature_name":["cases...
false	174,415.673	175,922	[{"feature_name":["cases...
false	174,415.673	177,520	[{"feature_name":["cases...
false	175,834.849	182,212	[{"feature_name":["cases...
false	175,834.849	186,297	[{"feature_name":["cases...
false	175,834.849	190,579	[{"feature_name":["cases...
false	173,401.506	194,288	[{"feature_name":["cases...
false	210,469.811	214,871	[{"feature_name":["cases...
false	214,166.061	217,666	[{"feature_name":["cases...
false	241,103.53	238,657	[{"feature_name":["cases...
false	274,732.626	275,513	[{"feature_name":["cases...
false	292,499.225	302,564	[{"feature_name":["cases...
false	333,752.614	351,667	[{"feature_name":["cases...
false	333,752.614	361,464	[{"feature_name":["cases...

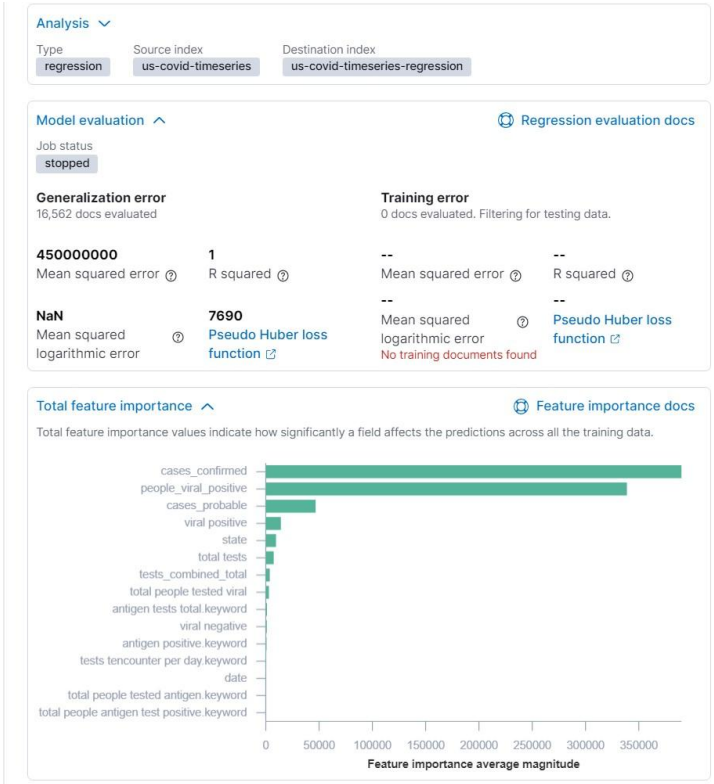

Creating Forecast

1. Once the data has been uploaded, go to Analytics > Machine Learning > Anomaly Detection > Jobs

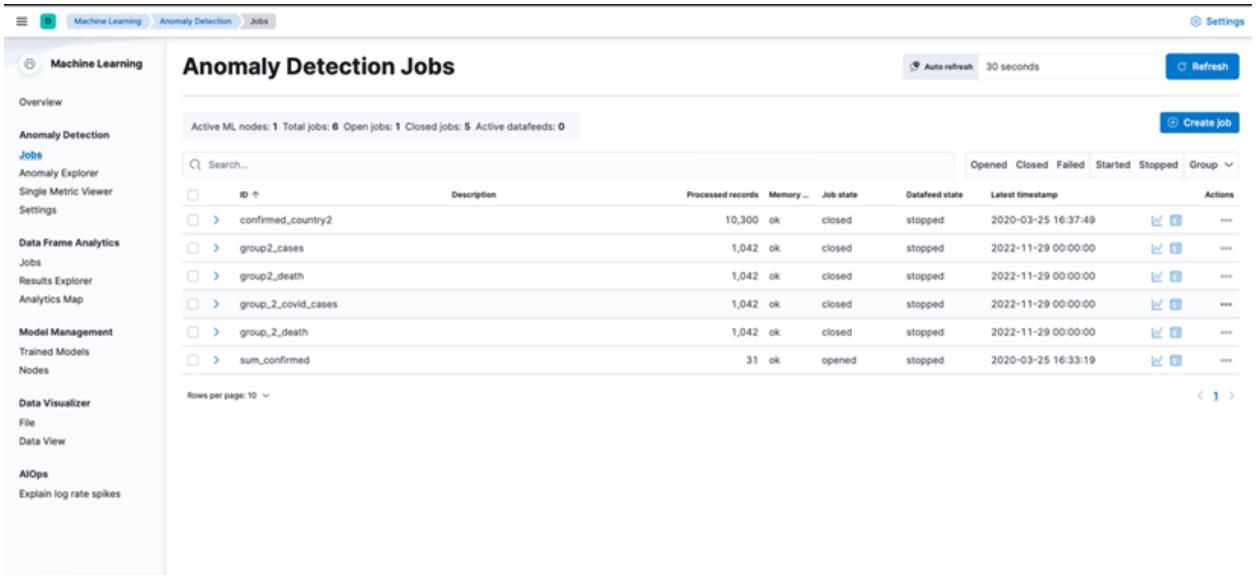

2. Now, press the **create job** button

3. Select the data that we have uploaded on Kibana (covid-index-county)

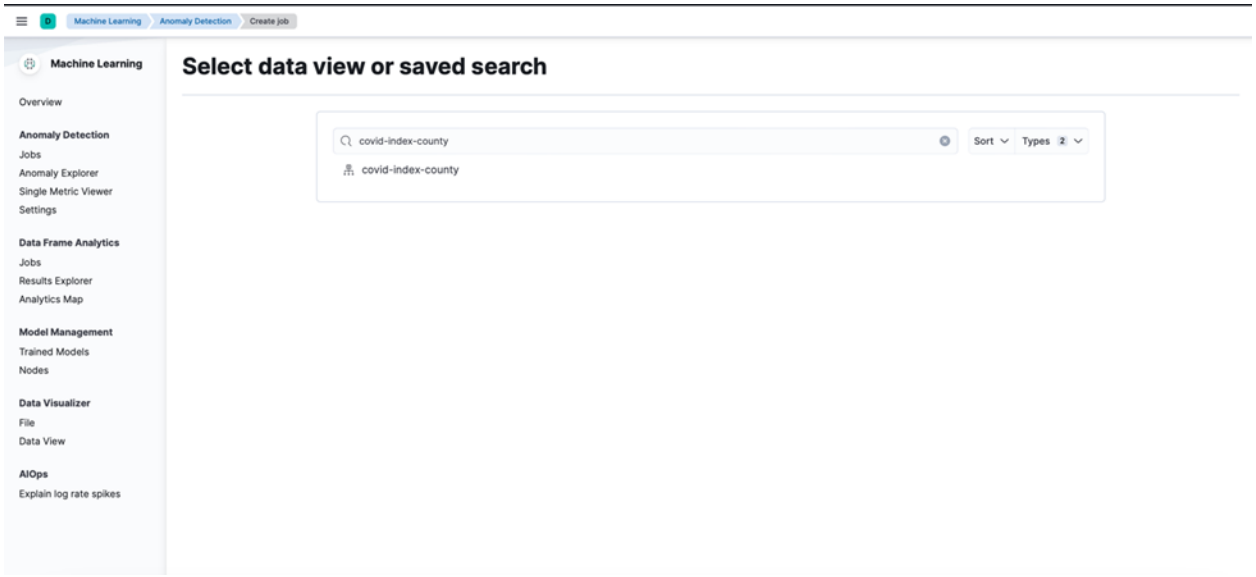

4. Once Selected, press on **single metrics**

5. On Create Job, Single Metrics, for the time range, press the '**Use full range data**' and press next

6. On the field, select **mean(case)** and press next

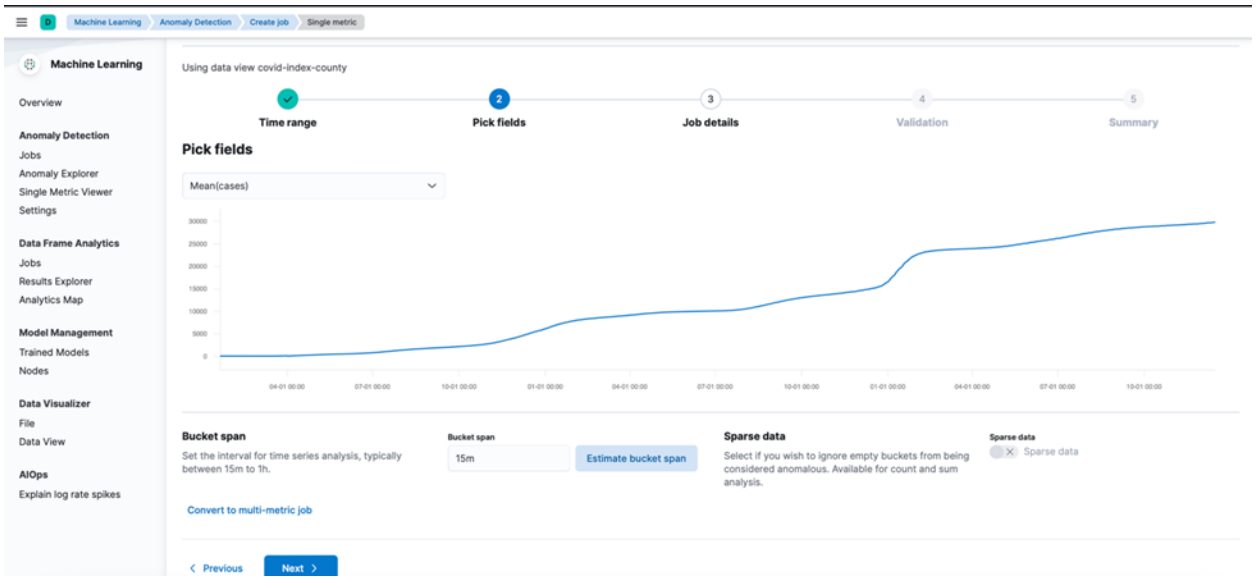

7. Enter the job id as 'covid-19-cases-forecast'

8. Press next repeatedly until you arrive on the summary page.

9. On the summary page, press the ‘create job’ button

10. Once the job has been successfully created, press the ‘view results button’

    a. On the results button, the data might not appear until the end because it is still loading. In this case, please wait a little and refresh the page until all the data appears.

11. Once the data is fully loaded press on the ‘forecast button’

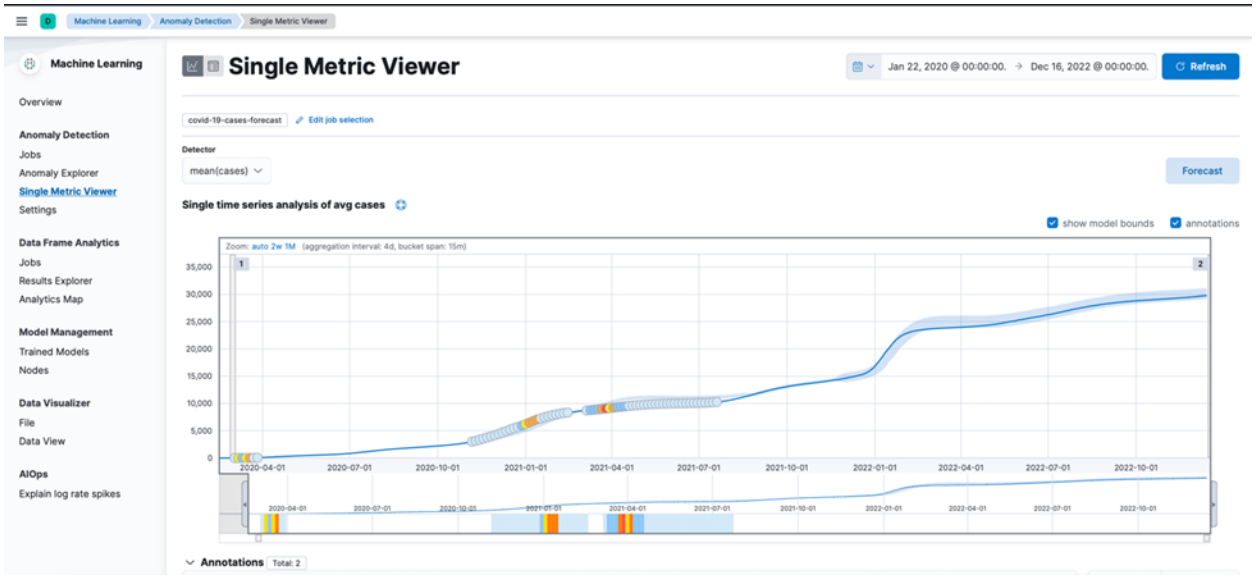

12. To create a forecast for a year, enter **365d** and **press forecast**

13. One completed, **drag** the **Single time series** to the end to the entire forecast for **next year**

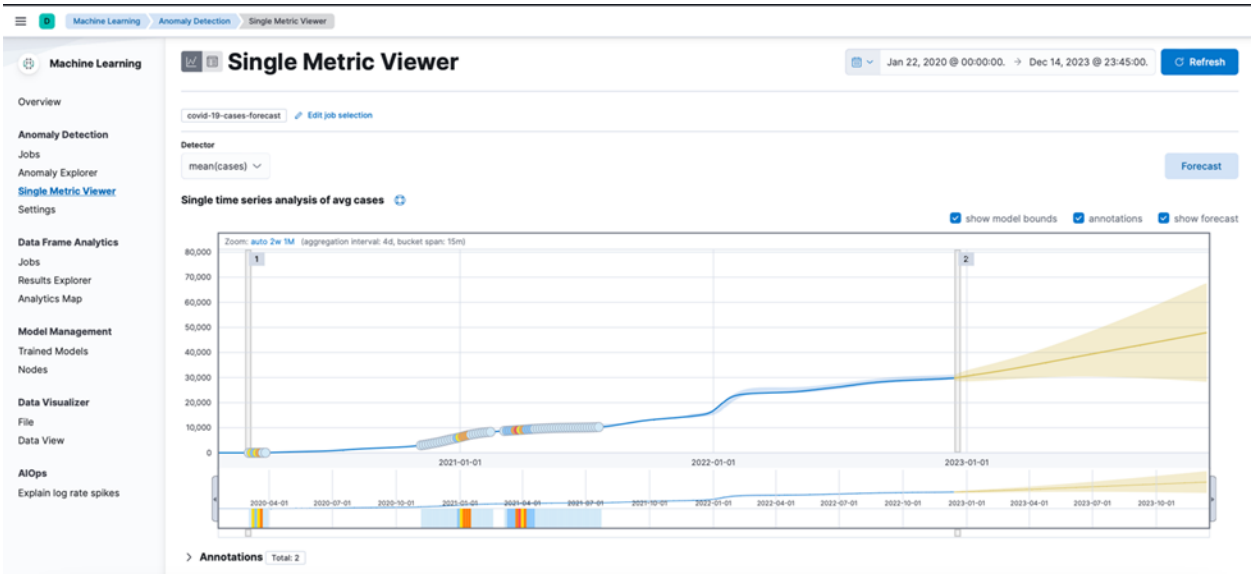

14. Now you have a forecast of the average cases for a year.

Future work, possibly, for your term project

You may increase the range of the forecast up to **3650 days** which is 10 years from now. Also you may change the average number of cases to average number of death when creating the job by changing the field from **mean(cases) to mean (death).**

Creating a Horizontal Bar & Vertical Stacked Bar Graph

Elasticsearch provides tools to create quick graphs. We will be using the tool “Lens” to create these 2 bar graphs.

1. First we will be heading to the Visualize Library by clicking on it under the Analytics

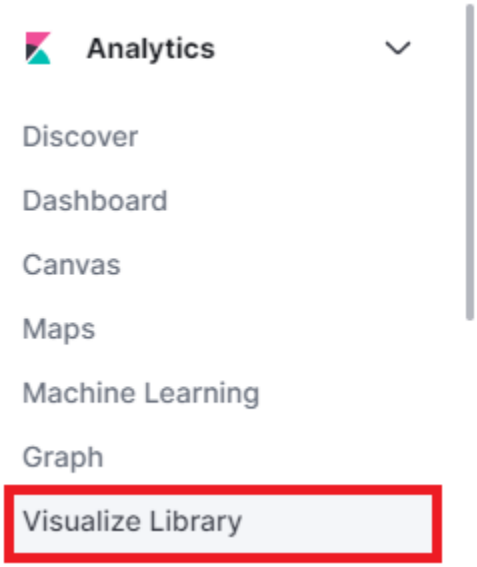

2. Let’s click on “**Create visualization**”.

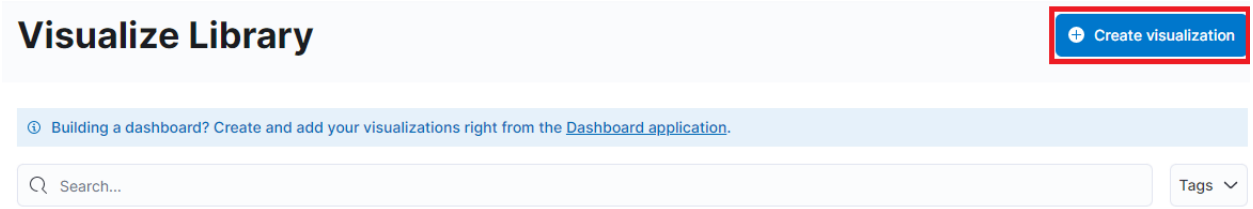

3. We will be prompted by several options to create a visualization, but we will go with Elastic’s easy, drag and drop, option: Lens.

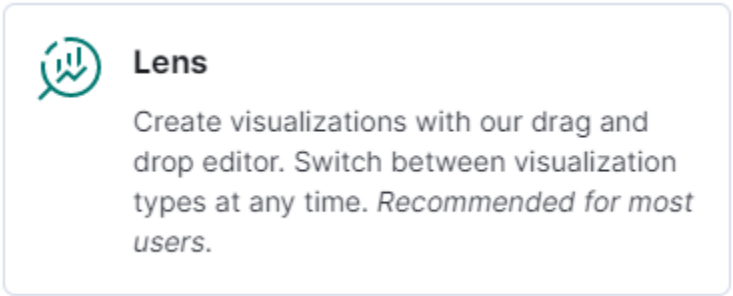

4. Let's make sure that we select "**Bar Horizontal**", since that's the one we are creating for now.

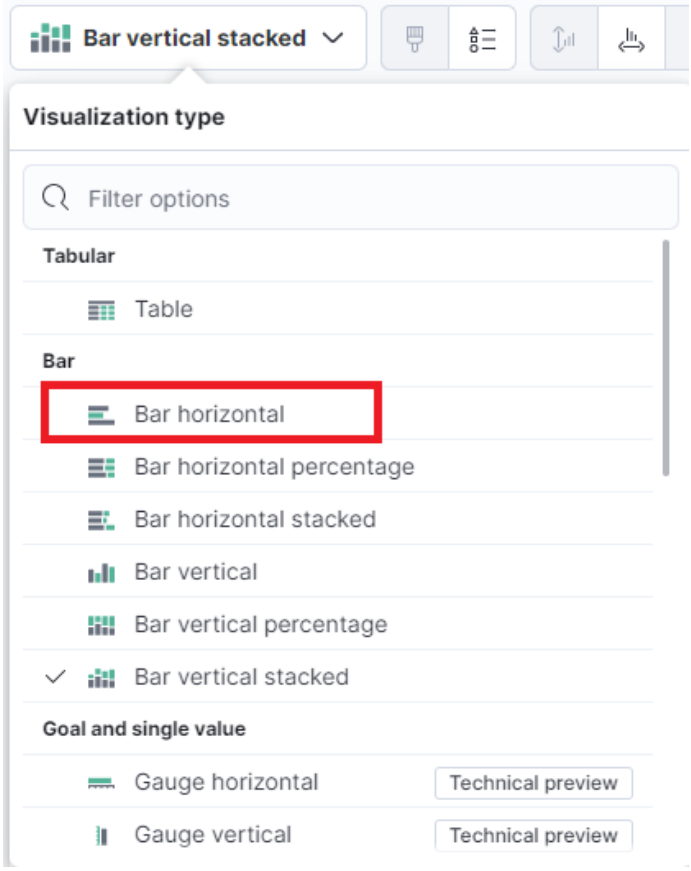

5. These are the options we are going with to see the Top 5 Counties with Confirmed cases in the U.S. Under functions, we are choosing "**Maximum**". Under field, we will be choosing "**cases**".

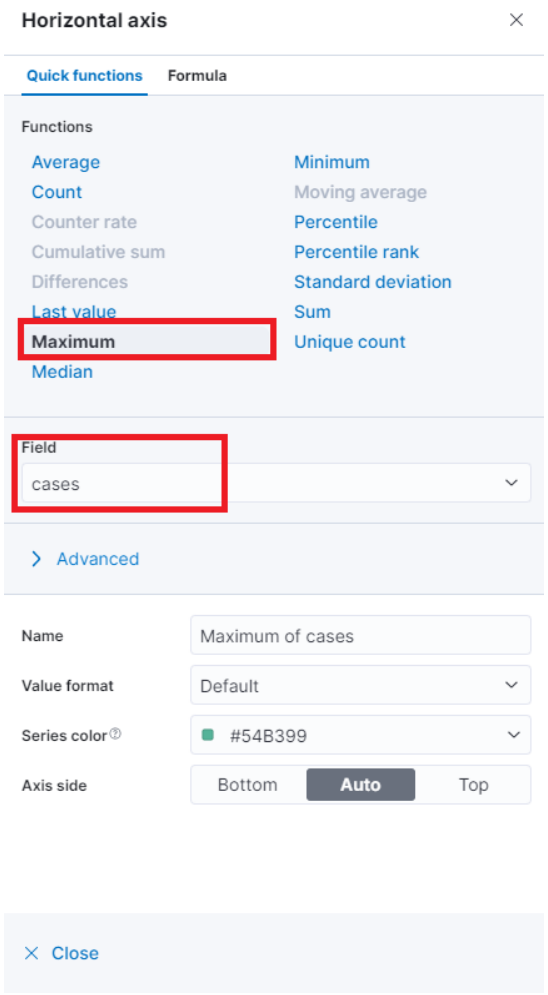

6. For the Vertical axis, we will be selecting “**Top Values**” for the Functions, “**County**” for fields, and Ranking by “**Maximum of cases**”.

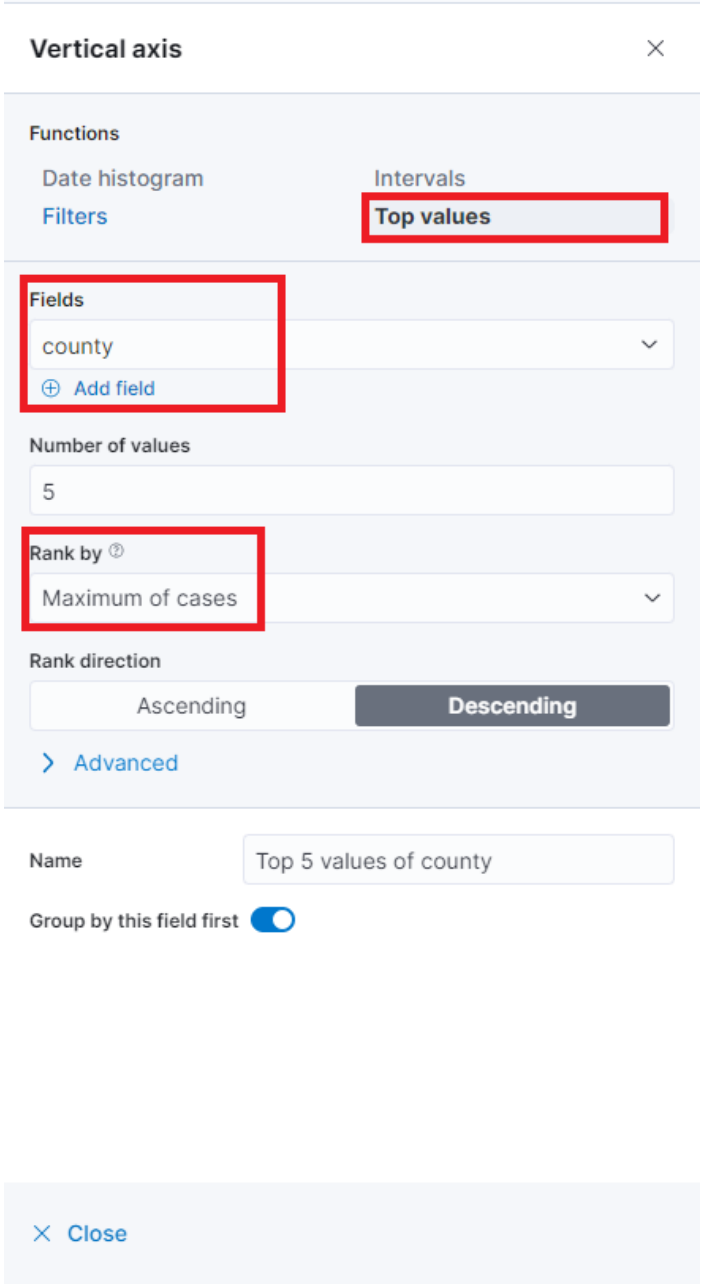

7. We want to break everything down by “**Top Values**” with “**county**” under Fields and Rank by “**Maximum Cases**”.

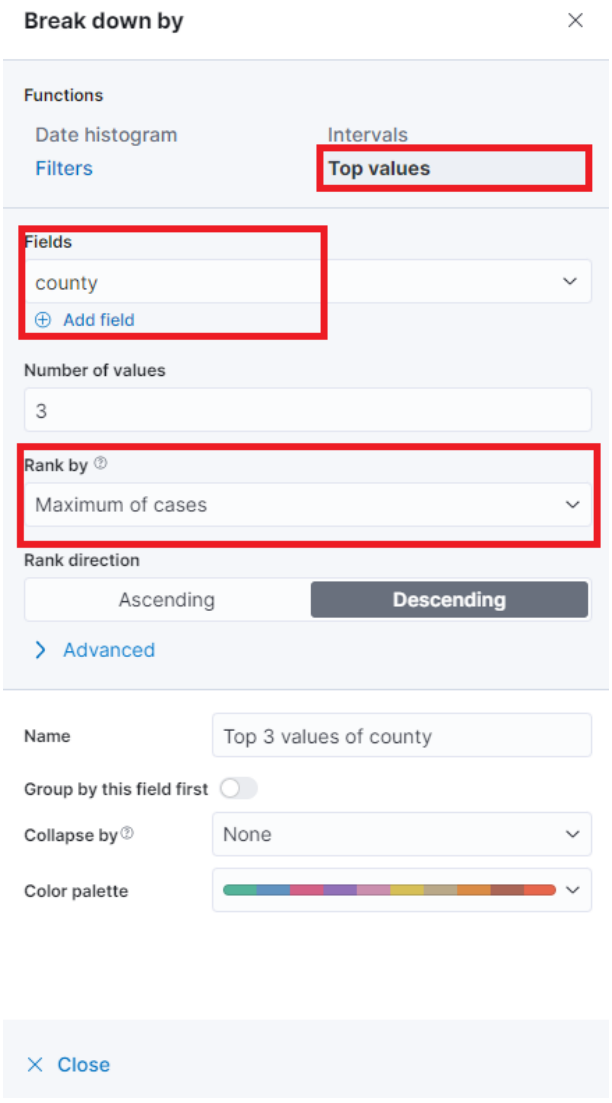

8. Below we can see our results. Los Angeles county seems to be the county with the highest number of cases.

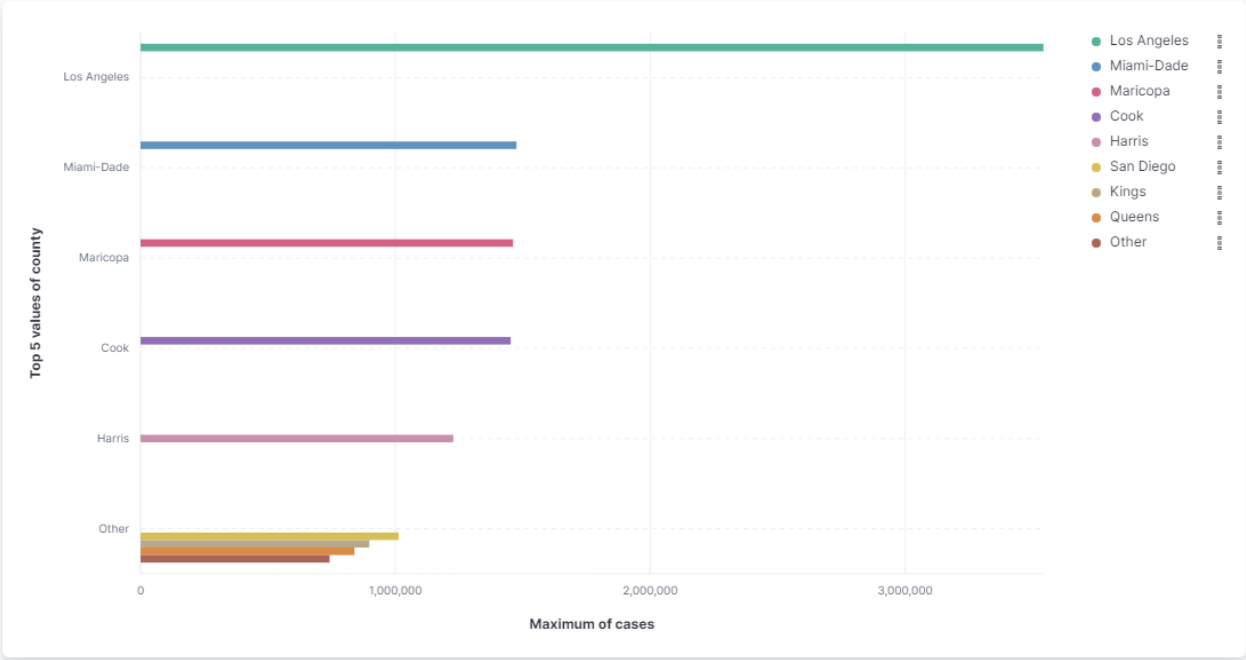

9. To create the Vertical Stacked Bar Graph, we're going to repeat the first initial steps. Up to the point where we let Elastic know what bar graph we want to create. Now we will choose "**Ber vertical stacked**".

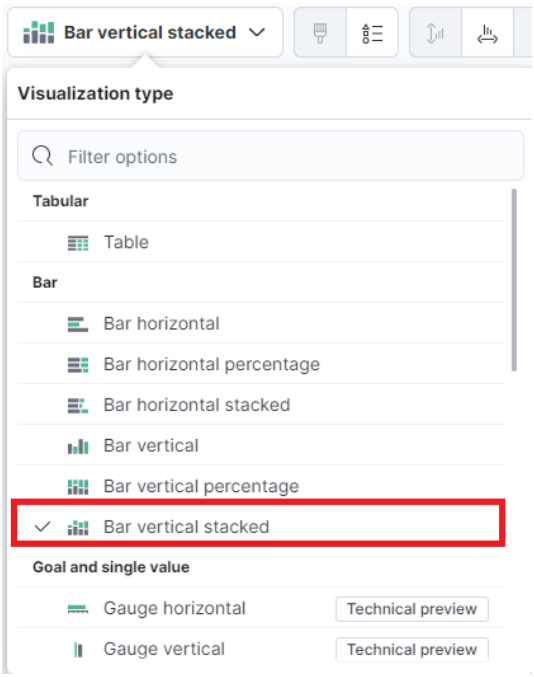

10. For our Vertical axis, "**Average**" will be our function and we will select **"deaths**" as our field.

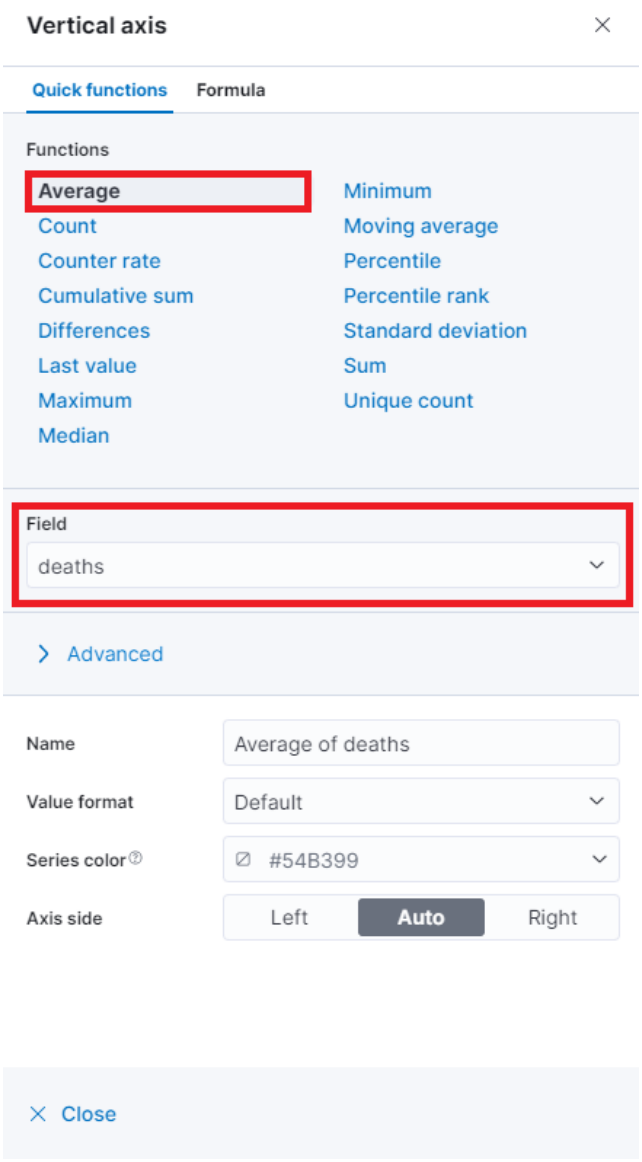

11. Our Vertical axes will have the “**Top Values**” as the Functions, with “**county**” selected as our Fields.

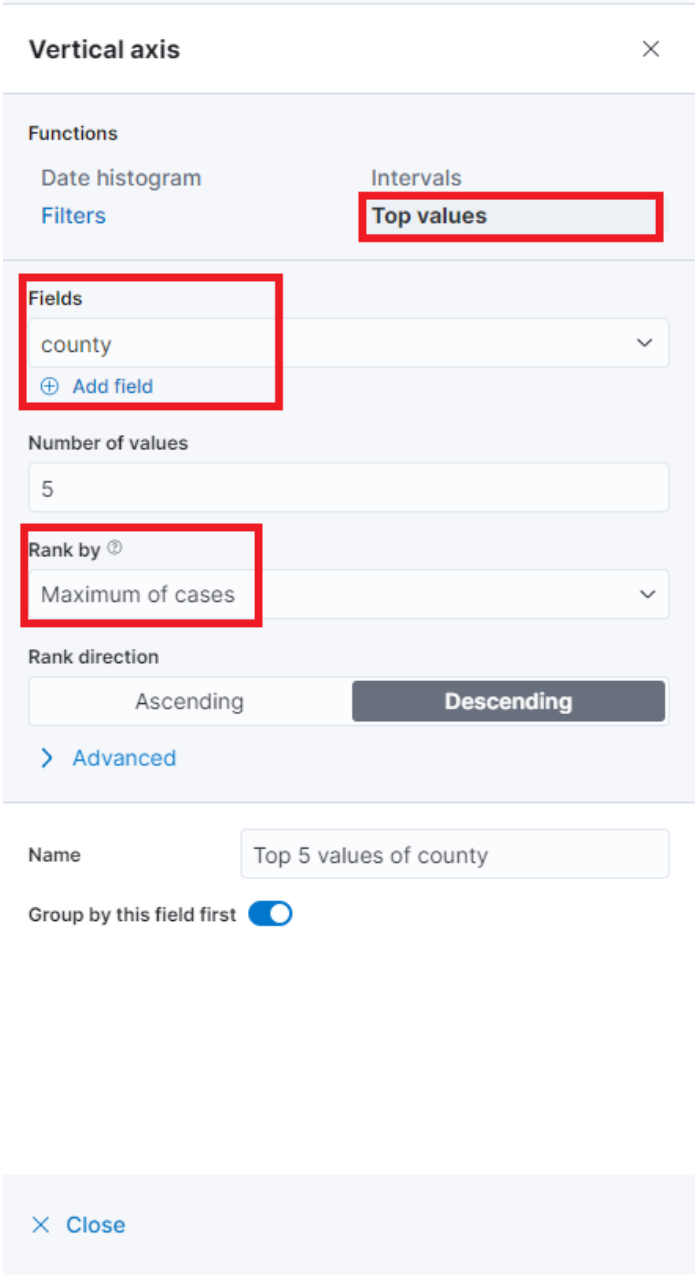

12. Let’s break down everything by “**Top Values**” under Functions, with “**state**” as our fields.

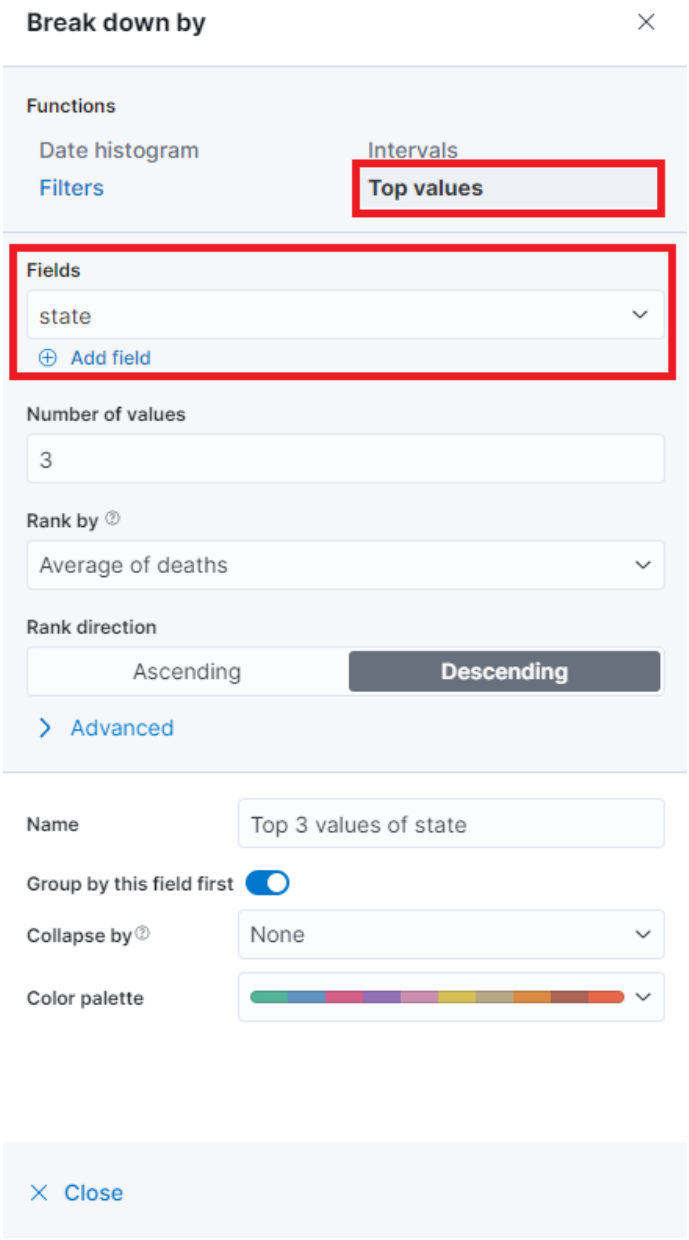

13. Once we have finished entering our settings, we can see our final results, the top 3 average amount of deaths per state.

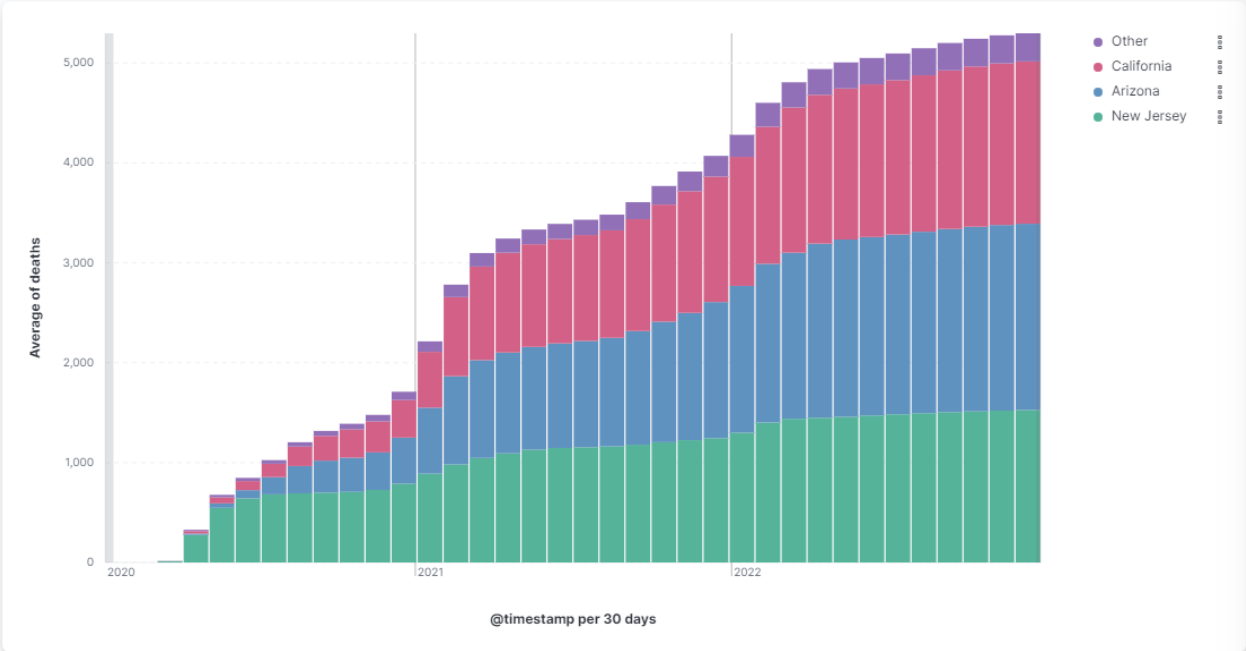

This is the end of the tutorial